

Formally Verified PCP Constructions

Serhat Emre Coban Davide Mazzali Khanh Nguyen
Vincent Palma Yanting Teng Thomas Vidick Yuxi Zheng

May 25, 2026

Chapter 1

BLR for general finite fields

1.1 BLR over General Finite Fields

Throughout this section, let $q = p^s$ for a prime p and $s \in \mathbb{N}$ with $s \geq 1$. All vectors are in $\mathbb{F}_q^n$, and all expectations are uniform over the indicated finite sets. Let $\omega_p = e^{2\pi i/p}$.

Definition 1.1.1 (Field trace). The field trace $\text{Tr} : \mathbb{F}_q \rightarrow \mathbb{F}_p$ is

$$\text{Tr}(z) := \sum_{i=0}^{s-1} z^{p^i}.$$

Definition 1.1.2 (Base additive character). The base additive character $\psi : \mathbb{F}_q \rightarrow \mathbb{C}$ is

$$\psi(z) := \omega_p^{\text{Tr}(z)}.$$

Definition 1.1.3 (Additive characters). For $\alpha \in \mathbb{F}_q^n$, define the additive character $\chi_\alpha : \mathbb{F}_q^n \rightarrow \mathbb{C}$ by

$$\chi_\alpha(x) := \psi(\langle \alpha, x \rangle) = \omega_p^{\text{Tr}(\langle \alpha, x \rangle)},$$

where $\langle \alpha, x \rangle = \sum_{i=1}^n \alpha_i x_i \in \mathbb{F}_q$.

Definition 1.1.4 (Expectation). For a function $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$, define its uniform expectation by

$$\mathbb{E}_{x \in \mathbb{F}_q^n}[h(x)] := \frac{1}{q^n} \sum_{x \in \mathbb{F}_q^n} h(x).$$

Definition 1.1.5 (Inner product). For functions $h_1, h_2 : \mathbb{F}_q^n \rightarrow \mathbb{C}$, define

$$\langle h_1, h_2 \rangle := \mathbb{E}_{x \in \mathbb{F}_q^n} [h_1(x) \overline{h_2(x)}].$$

Definition 1.1.6 (L_2 norm). For a function $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$, define its L_2 norm by

$$\|h\|_2 := \left(\mathbb{E}_{x \in \mathbb{F}_q^n} [|h(x)|^2] \right)^{1/2} = \sqrt{\langle h, h \rangle}.$$

Definition 1.1.7 (Fourier coefficient). For $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$ and $\alpha \in \mathbb{F}_q^n$, define

$$\hat{h}(\alpha) := \langle h, \chi_\alpha \rangle = \mathbb{E}_{x \in \mathbb{F}_q^n} [h(x) \overline{\chi_\alpha(x)}].$$

Lemma 1.1.8 (Characters are orthonormal). *For every $\alpha, \beta \in \mathbb{F}_q^n$,*

$$\langle \chi_\alpha, \chi_\beta \rangle = \begin{cases} 1 & \text{if } \alpha = \beta, \\ 0 & \text{if } \alpha \neq \beta. \end{cases}$$

Proof. First observe that for every $\alpha, \beta \in \mathbb{F}_q^n$ and every $x \in \mathbb{F}_q^n$,

$$\chi_\alpha(x) \overline{\chi_\beta(x)} = \omega_p^{\text{Tr}(\langle \alpha, x \rangle)} \omega_p^{-\text{Tr}(\langle \beta, x \rangle)} = \omega_p^{\text{Tr}(\langle \alpha - \beta, x \rangle)} = \chi_{\alpha - \beta}(x),$$

where we used the $\mathbb{F}_p$ -linearity of the trace map.

If $\alpha = \beta$, then $\alpha - \beta = 0$, so $\chi_{\alpha - \beta}(x) = 1$ for every $x \in \mathbb{F}_q^n$. Hence

$$\mathbb{E}_{x \in \mathbb{F}_q^n} [\chi_\alpha(x) \overline{\chi_\beta(x)}] = \mathbb{E}_{x \in \mathbb{F}_q^n} [1] = 1.$$

Now suppose $\alpha \neq \beta$. Let $\gamma := \alpha - \beta$. Then $\gamma \neq 0$, so there exists an index j such that $\gamma_j \neq 0$. Fix all coordinates of x except x_j . For the fixed remaining coordinates, the map

$$x_j \mapsto \langle \gamma, x \rangle$$

has the form

$$x_j \mapsto \gamma_j x_j + c$$

for some constant $c \in \mathbb{F}_q$. Since $\gamma_j \neq 0$, multiplication by γ_j is a bijection on $\mathbb{F}_q$, and therefore the affine map $x_j \mapsto \gamma_j x_j + c$ is also a bijection on $\mathbb{F}_q$. Thus, as x_j ranges uniformly over $\mathbb{F}_q$, the value $\langle \gamma, x \rangle$ also ranges uniformly over $\mathbb{F}_q$. Averaging first over x_j and then over the remaining coordinates gives

$$\mathbb{E}_{x \in \mathbb{F}_q^n} [\chi_\gamma(x)] = \mathbb{E}_{t \in \mathbb{F}_q} [\omega_p^{\text{Tr}(t)}].$$

It remains to show that this last average is zero. The trace map $\text{Tr} : \mathbb{F}_q \rightarrow \mathbb{F}_p$ is a nonzero $\mathbb{F}_p$ -linear map, so every fiber of Tr has the same cardinality. Since $|\mathbb{F}_q| = q$ and $|\mathbb{F}_p| = p$, each fiber has cardinality q/p . Therefore

$$\begin{aligned} \mathbb{E}_{t \in \mathbb{F}_q} [\omega_p^{\text{Tr}(t)}] &= \frac{1}{q} \sum_{t \in \mathbb{F}_q} \omega_p^{\text{Tr}(t)} \\ &= \frac{1}{q} \sum_{r \in \mathbb{F}_p} \sum_{\substack{t \in \mathbb{F}_q \\ \text{Tr}(t) = r}} \omega_p^r \\ &= \frac{1}{q} \sum_{r \in \mathbb{F}_p} \frac{q}{p} \omega_p^r \\ &= \frac{1}{p} \sum_{r \in \mathbb{F}_p} \omega_p^r \\ &= 0, \end{aligned}$$

because the sum of all p -th roots of unity is zero. Hence, if $\alpha \neq \beta$, then

$$\mathbb{E}_{x \in \mathbb{F}_q^n} [\chi_\alpha(x) \overline{\chi_\beta(x)}] = 0.$$

Combining the two cases proves the stated orthogonality relation. The equivalent formulation follows by taking $\beta = 0$. $\square$

Lemma 1.1.9 (Fourier inversion). *For every function $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$ we have*

$$h(x) = \sum_{\alpha \in \mathbb{F}_q^n} \hat{h}(\alpha) \chi_\alpha(x) \quad \text{for every } x \in \mathbb{F}_q^n.$$

Proof. By Theorem 1.1.8, the additive characters

$$\{\chi_\alpha : \alpha \in \mathbb{F}_q^n\}$$

are orthonormal with respect to the inner product

$$\langle h_1, h_2 \rangle = \mathbb{E}_{x \in \mathbb{F}_q^n} [h_1(x) \overline{h_2(x)}].$$

Moreover, these characters span the vector space of all complex-valued functions on $\mathbb{F}_q^n$. Indeed, there are exactly q^n characters, and the space of functions $\mathbb{F}_q^n \rightarrow \mathbb{C}$ has dimension q^n . Since the characters are nonzero and pairwise orthogonal, they are linearly independent. Hence they form an orthonormal basis.

Therefore every function $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$ has a unique expansion in this basis:

$$h = \sum_{\alpha \in \mathbb{F}_q^n} c_\alpha \chi_\alpha$$

for some coefficients $c_\alpha \in \mathbb{C}$. Taking the inner product of both sides with χ_β , we obtain

$$\langle h, \chi_\beta \rangle = \sum_{\alpha \in \mathbb{F}_q^n} c_\alpha \langle \chi_\alpha, \chi_\beta \rangle.$$

By orthonormality, all terms in the sum vanish except the term $\alpha = \beta$, and $\langle \chi_\beta, \chi_\beta \rangle = 1$. Hence

$$\langle h, \chi_\beta \rangle = c_\beta.$$

By definition of the Fourier coefficient,

$$\hat{h}(\beta) = \langle h, \chi_\beta \rangle.$$

Thus $c_\beta = \hat{h}(\beta)$ for every $\beta \in \mathbb{F}_q^n$. Therefore

$$h = \sum_{\alpha \in \mathbb{F}_q^n} \hat{h}(\alpha) \chi_\alpha.$$

Evaluating this identity at $x \in \mathbb{F}_q^n$ gives the Fourier inversion formula

$$h(x) = \sum_{\alpha \in \mathbb{F}_q^n} \hat{h}(\alpha) \chi_\alpha(x).$$

□

Lemma 1.1.10 (Plancherel identity). *For every pair of functions $f, g : \mathbb{F}_q^n \rightarrow \mathbb{C}$,*

$$\sum_{\alpha \in \mathbb{F}_q^n} \hat{f}(\alpha) \overline{\hat{g}(\alpha)} = \langle f, g \rangle.$$

Proof. By Theorem 1.1.9,

$$f = \sum_{\alpha \in \mathbb{F}_q^n} \hat{f}(\alpha) \chi_\alpha.$$

Taking the inner product with g and using linearity in the first argument gives

$$\langle f, g \rangle = \sum_{\alpha \in \mathbb{F}_q^n} \hat{f}(\alpha) \langle \chi_\alpha, g \rangle.$$

Since

$$\langle \chi_\alpha, g \rangle = \overline{\langle g, \chi_\alpha \rangle} = \overline{\hat{g}(\alpha)},$$

we obtain

$$\langle f, g \rangle = \sum_{\alpha \in \mathbb{F}_q^n} \hat{f}(\alpha) \overline{\hat{g}(\alpha)}.$$

□

Lemma 1.1.11 (Parseval identity). *For every function $h : \mathbb{F}_q^n \rightarrow \mathbb{C}$,*

$$\sum_{\alpha \in \mathbb{F}_q^n} |\hat{h}(\alpha)|^2 = \|h\|_2^2.$$

Proof. Apply Theorem 1.1.10 with $f = g = h$:

$$\sum_{\alpha \in \mathbb{F}_q^n} \hat{h}(\alpha) \overline{\hat{h}(\alpha)} = \langle h, h \rangle.$$

The left-hand side is

$$\sum_{\alpha \in \mathbb{F}_q^n} |\hat{h}(\alpha)|^2,$$

and by the definitions of the inner product and the L_2 norm, the right-hand side is

$$\langle h, h \rangle = \mathbb{E}_{x \in \mathbb{F}_q^n} [|h(x)|^2] = \|h\|_2^2.$$

□

Definition 1.1.12 (Nonzero field elements). Write

$$\mathbb{F}_q^\times := \{c \in \mathbb{F}_q : c \neq 0\}.$$

Definition 1.1.13 (Fourier expansion of a finite-field-valued function). Let $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$. For each $c \in \mathbb{F}_q^\times$, define the c -phase lift

$$\varphi_c(x) := \omega_p^{\text{Tr}(cf(x))}.$$

The Fourier expansion of f is the set $\{\varphi_c\}_{c \in \mathbb{F}_q^\times}$.

Lemma 1.1.14 (Character-sum indicator). *For every $u, v \in \mathbb{F}_q$,*

$$\mathbb{1}[u = v] = \frac{1}{q} \sum_{c \in \mathbb{F}_q} \omega_p^{\text{Tr}(c(u-v))}.$$

Proof. For each $z \in \mathbb{F}_q$ we define

$$\psi(z) := \omega_p^{\text{Tr}(z)}$$

be the standard additive character of $\mathbb{F}_q$. We prove that, for every $t \in \mathbb{F}_q$,

$$\mathbb{1}[t = 0] = \frac{1}{q} \sum_{c \in \mathbb{F}_q} \psi(ct).$$

The desired statement follows by taking $t = u - v$.

If $t = 0$, then $ct = 0$ for every $c \in \mathbb{F}_q$, so

$$\frac{1}{q} \sum_{c \in \mathbb{F}_q} \psi(ct) = \frac{1}{q} \sum_{c \in \mathbb{F}_q} 1 = 1.$$

Now suppose $t \neq 0$. Then multiplication by t is a bijection $\mathbb{F}_q \rightarrow \mathbb{F}_q$, so

$$\sum_{c \in \mathbb{F}_q} \psi(ct) = \sum_{z \in \mathbb{F}_q} \psi(z).$$

We claim that the latter sum is 0. Let

$$S := \sum_{z \in \mathbb{F}_q} \psi(z).$$

Since $\text{Tr} : \mathbb{F}_q \rightarrow \mathbb{F}_p$ is a nonzero $\mathbb{F}_p$ -linear map, it is surjective. Hence there exists $a \in \mathbb{F}_q$ such that $\text{Tr}(a) = 1$. Using the bijection $z \mapsto z + a$, we get

$$S = \sum_{z \in \mathbb{F}_q} \psi(z + a) = \sum_{z \in \mathbb{F}_q} \omega_p^{\text{Tr}(z+a)} = \sum_{z \in \mathbb{F}_q} \omega_p^{\text{Tr}(z)} \omega_p^{\text{Tr}(a)} = \omega_p S.$$

Since $\omega_p \neq 1$, this implies $S = 0$. Therefore, when $t \neq 0$,

$$\frac{1}{q} \sum_{c \in \mathbb{F}_q} \psi(ct) = 0.$$

Combining the two cases gives

$$\mathbb{1}[t = 0] = \frac{1}{q} \sum_{c \in \mathbb{F}_q} \omega_p^{\text{Tr}(ct)}.$$

Substituting $t = u - v$, we obtain

$$\mathbb{1}[u = v] = \frac{1}{q} \sum_{c \in \mathbb{F}_q} \omega_p^{\text{Tr}(c(u-v))}.$$

□

Definition 1.1.15 (Relative Hamming distance). For $f, g : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, define

$$\delta(f, g) := \Pr_{x \in \mathbb{F}_q^n} [f(x) \neq g(x)].$$

Definition 1.1.16 (Linear functions). For $\alpha \in \mathbb{F}_q^n$, define the linear scalar function $\ell_\alpha : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ by

$$\ell_\alpha(x) := \langle \alpha, x \rangle.$$

A function $h : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ is linear if $h = \ell_\alpha$ for some $\alpha \in \mathbb{F}_q^n$. The finite set of all linear scalar functions is

$$\mathcal{LIN} := \{\ell_\alpha \mid \alpha \in \mathbb{F}_q^n\}.$$

Definition 1.1.17 (Distance to linearity). The distance from $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ to linearity is

$$\delta(f, \mathcal{LIN}) := \min_{g \in \mathcal{LIN}} \delta(f, g).$$

Lemma 1.1.18 (Well-separatedness of $\mathcal{LIN}$). For any distinct $f, g \in \mathcal{LIN}$, one has $\delta(f, g) = 1 - 1/|\mathbb{F}|$.

Proof. Let $a, b \in \mathbb{F}^n$ be the unique vectors such that $f(x) = \langle a, x \rangle$ and $g(x) = \langle b, x \rangle$ for each x . Let $c = a - b$, which must be a non-zero vector. In particular, there exists $i \in [n]$ such that $c_i \neq 0$. Take any permutation $\sigma : [n] \rightarrow [n]$ such that $\sigma(i) = n$. Then

$$\begin{aligned} \delta(f, g) &= 1 - \Pr_{x \sim \mathbb{F}^n} [\langle c, x \rangle = 0] \\ &= 1 - \frac{1}{|\mathbb{F}|^{n-1}} \sum_{x_1, \dots, x_{n-1} \in \mathbb{F}} \frac{1}{|\mathbb{F}|} \sum_{x_n \in \mathbb{F}} \mathbb{1}[c_i \cdot x_n = - \sum_{j \in [n] \setminus i} c_j \cdot x_{\sigma(j)}] \\ &= 1 - \frac{1}{|\mathbb{F}|^{n-1}} \sum_{x_1, \dots, x_{n-1} \in \mathbb{F}} \frac{1}{|\mathbb{F}|} \sum_{x_n \in \mathbb{F}} \mathbb{1}[x_n = -c_i^{-1} \cdot \sum_{j \in [n] \setminus i} c_j \cdot x_{\sigma(j)}] \end{aligned}$$

where the last equality follows because $c_i \neq 0$. Now note that $\mathbb{1}[x_n = -c_i^{-1} \cdot \sum_{j \in [n] \setminus i} c_j \cdot x_{\sigma(j)}]$ evaluates to 1 for exactly one value of x_n . Thus we continue from above and get

$$\delta(f, g) = 1 - \frac{1}{|\mathbb{F}|^{n-1}} \sum_{x_1, \dots, x_{n-1} \in \mathbb{F}} \frac{1}{|\mathbb{F}|} = 1 - \frac{1}{|\mathbb{F}|}.$$

□

Lemma 1.1.19 (Distance formula). Let $f, g : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$. Let $\{\varphi_c\}_{c \in \mathbb{F}_q^\times}$ and $\{\gamma_c\}_{c \in \mathbb{F}_q^\times}$ be their Fourier expansions. Then

$$\delta(f, g) = 1 - \frac{1}{q} \left(1 + \sum_{c \in \mathbb{F}_q^\times} \langle \varphi_c, \gamma_c \rangle \right).$$

Proof. By Theorem 1.1.14,

$$\Pr_x[f(x) = g(x)] = \mathbb{E}_x \left[\frac{1}{q} \sum_{c \in \mathbb{F}_q} \omega_p^{\text{Tr}(c(f(x) - g(x)))} \right].$$

The term $c = 0$ contributes $1/q$. For $c \neq 0$,

$$\omega_p^{\text{Tr}(c(f(x) - g(x)))} = \varphi_c(x) \overline{\gamma_c(x)}.$$

Therefore

$$\Pr_x[f(x) = g(x)] = \frac{1}{q} \left(1 + \sum_{c \in \mathbb{F}_q^\times} \langle \varphi_c, \gamma_c \rangle \right),$$

which gives the first formula after using $\delta(f, g) = 1 - \Pr_x[f(x) = g(x)]$. □

Lemma 1.1.20 (Fourier expansion of a linear function). *Fix $\alpha \in \mathbb{F}_q^n$ and let $\{\lambda_c\}_{c \in \mathbb{F}_q^\times}$ be the Fourier expansion of ℓ_α . Then*

$$\lambda_c = \chi_{c\alpha}.$$

Proof. For every $x \in \mathbb{F}_q^n$,

$$\lambda_c(x) = \omega_p^{\text{Tr}(c\langle\alpha, x\rangle)} = \omega_p^{\text{Tr}(\langle c\alpha, x\rangle)} = \chi_{c\alpha}(x).$$

□

Definition 1.1.21 (Linear Fourier score). For $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ and $\alpha \in \mathbb{F}_q^n$, define the linear Fourier score of f at α by

$$S_\alpha(f) := \sum_{c \in \mathbb{F}_q^\times} \widehat{\varphi}_c(c\alpha),$$

where $\{\varphi_c\}_{c \in \mathbb{F}_q^\times}$ is the Fourier expansion of f .

Lemma 1.1.22 (Distance to linearity in Fourier form). *Let $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ have Fourier expansion $\{\varphi_c\}_{c \in \mathbb{F}_q^\times}$. Then*

$$\delta(f, \mathcal{LIN}) = 1 - \frac{1}{q} \left(1 + \max_{\alpha \in \mathbb{F}_q^n} S_\alpha(f) \right).$$

Moreover, for every $\alpha \in \mathbb{F}_q^n$, the score $S_\alpha(f)$ is real and satisfies $-1 \leq S_\alpha(f) \leq q-1$.

Proof. Apply Theorem 1.1.19 with $g = \ell_\alpha$ and use Theorem 1.1.20. This gives

$$\delta(f, \ell_\alpha) = 1 - \frac{1}{q} (1 + S_\alpha(f)).$$

Taking the minimum over α gives the claimed formula for $\delta(f, \mathcal{LIN})$. For fixed α , the same formula expresses $S_\alpha(f)$ as $q(1 - \delta(f, \ell_\alpha)) - 1$. Hence $S_\alpha(f)$ is real, and since $0 \leq \delta(f, \ell_\alpha) \leq 1$, we get $-1 \leq S_\alpha(f) \leq q-1$. □

Definition 1.1.23 (Finite-field BLR test). Given oracle access to $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, the finite-field BLR verifier V_{BLR}^f samples $a, b \in \mathbb{F}_q^\times$ and $x, y \in \mathbb{F}_q^n$ independently and uniformly at random, and accepts iff

$$af(x) + bf(y) = f(ax + by).$$

That is,

$$V_{\text{BLR}}^f = 1 \iff af(x) + bf(y) = f(ax + by).$$

Lemma 1.1.24 (Completeness). *If $f \in \mathcal{LIN}$, then*

$$\Pr[V_{\text{BLR}}^f = 1] = 1.$$

Proof. If $f = \ell_\alpha$ for some $\alpha \in \mathbb{F}_q^n$, then for every $a, b \in \mathbb{F}_q^\times$ and $x, y \in \mathbb{F}_q^n$,

$$f(ax + by) = \langle \alpha, ax + by \rangle = a\langle \alpha, x \rangle + b\langle \alpha, y \rangle = af(x) + bf(y).$$

Thus the verifier accepts for every choice of randomness. □

Lemma 1.1.25 (Exact BLR acceptance formula). *Let $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$. Then*

$$\Pr[V_{\text{BLR}}^f = 1] = \frac{1}{q} \left(1 + \frac{1}{(q-1)^2} \sum_{\alpha \in \mathbb{F}_q^n} S_\alpha(f)^3 \right).$$

Proof. By Theorem 1.1.14, for fixed a, b, x, y ,

$$\mathbb{1}[af(x) + bf(y) = f(ax + by)] = \frac{1}{q} \sum_{c \in \mathbb{F}_q} \omega_p^{\text{Tr}(c(af(x) + bf(y) - f(ax + by)))}.$$

Taking expectation gives

$$\Pr[V_{\text{BLR}}^f = 1] = \frac{1}{q} + \frac{1}{q} \mathbb{E}_{a,b,x,y} \left[\sum_{c \in \mathbb{F}_q^\times} \varphi_{ca}(x) \varphi_{cb}(y) \varphi_{-c}(ax + by) \right].$$

Expand the three complex-valued functions into their Fourier expansions:

$$\begin{aligned} & \mathbb{E}_{a,b,x,y} \left[\sum_{c \in \mathbb{F}_q^\times} \varphi_{ca}(x) \varphi_{cb}(y) \varphi_{-c}(ax + by) \right] \\ &= \mathbb{E}_{a,b} \left[\sum_{c \in \mathbb{F}_q^\times} \sum_{\alpha, \beta, \gamma \in \mathbb{F}_q^n} \widehat{\varphi_{ca}}(\alpha) \widehat{\varphi_{cb}}(\beta) \widehat{\varphi_{-c}}(\gamma) \mathbb{E}_{x,y} [\chi_\alpha(x) \chi_\beta(y) \chi_\gamma(ax + by)] \right]. \end{aligned}$$

Since

$$\chi_\gamma(ax + by) = \chi_{a\gamma}(x) \chi_{b\gamma}(y),$$

Theorem 1.1.8 implies that the inner expectation is 1 exactly when

$$\alpha + a\gamma = 0 \quad \text{and} \quad \beta + b\gamma = 0,$$

and is 0 otherwise. Therefore the last display equals

$$\mathbb{E}_{a,b} \left[\sum_{c \in \mathbb{F}_q^\times} \sum_{\alpha \in \mathbb{F}_q^n} \widehat{\varphi_{ca}}(\alpha) \widehat{\varphi_{cb}}(a^{-1}b\alpha) \widehat{\varphi_{-c}}(-a^{-1}\alpha) \right].$$

Substitute $\alpha = ca\eta$. Then this becomes

$$\mathbb{E}_{a,b} \left[\sum_{c \in \mathbb{F}_q^\times} \sum_{\eta \in \mathbb{F}_q^n} \widehat{\varphi_{ca}}(ca\eta) \widehat{\varphi_{cb}}(cb\eta) \widehat{\varphi_{-c}}(-c\eta) \right].$$

As a, b, c range over $\mathbb{F}_q^\times$, so do ca, cb , and $-c$. Hence this equals

$$\frac{1}{(q-1)^2} \sum_{\eta \in \mathbb{F}_q^n} \left(\sum_{d \in \mathbb{F}_q^\times} \widehat{\varphi_d}(d\eta) \right)^3 = \frac{1}{(q-1)^2} \sum_{\eta \in \mathbb{F}_q^n} S_\eta(f)^3.$$

Substituting into the expression for $\Pr[V_{\text{BLR}}^f = 1]$ proves the claim. $\square$

Lemma 1.1.26 (Second moment bound). *For every $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, one has*

$$\sum_{\alpha \in \mathbb{F}_q^n} S_\alpha(f)^2 \leq (q-1)^2.$$

Proof. Expanding the square gives

$$\sum_{\alpha \in \mathbb{F}_q^n} S_\alpha(f)^2 = \sum_{a, b \in \mathbb{F}_q^\times} \sum_{\alpha \in \mathbb{F}_q^n} \widehat{\varphi}_a(a\alpha) \widehat{\varphi}_b(b\alpha).$$

For fixed $a, b \in \mathbb{F}_q^\times$, Cauchy-Schwarz and Theorem 1.1.11 give

$$\begin{aligned} \left| \sum_{\alpha \in \mathbb{F}_q^n} \widehat{\varphi}_a(a\alpha) \widehat{\varphi}_b(b\alpha) \right| &\leq \left(\sum_{\alpha \in \mathbb{F}_q^n} |\widehat{\varphi}_a(a\alpha)|^2 \right)^{1/2} \left(\sum_{\alpha \in \mathbb{F}_q^n} |\widehat{\varphi}_b(b\alpha)|^2 \right)^{1/2} \\ &= 1. \end{aligned}$$

The equality uses that multiplication by a and by b permutes $\mathbb{F}_q^n$, and that Parseval plus $|\varphi_a(x)| = |\varphi_b(x)| = 1$ for every x make both squared L_2 norms equal to 1. Summing over the $(q-1)^2$ choices of (a, b) gives the result. $\square$

Theorem 1.1.27 (Soundness of the finite-field BLR test). *For every function $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, one has*

$$\Pr[\mathbf{V}_{\text{BLR}}^f = 1] \leq 1 - \delta(f, \mathcal{LIN}).$$

Proof. Let $S_\alpha(f)$ be as in Theorem 1.1.21, and let

$$M := \max_{\alpha \in \mathbb{F}_q^n} S_\alpha(f).$$

By Theorem 1.1.22,

$$1 - \delta(f, \mathcal{LIN}) = \frac{1}{q}(1 + M).$$

By Theorem 1.1.22, each $S_\alpha(f)$ is real and $S_\alpha(f) \leq M$. Therefore $S_\alpha(f)^3 \leq M S_\alpha(f)^2$ for every α , because $S_\alpha(f)^2 \geq 0$. Using Theorem 1.1.25 and Theorem 1.1.26,

$$\begin{aligned} \Pr[\mathbf{V}_{\text{BLR}}^f = 1] &= \frac{1}{q} \left(1 + \frac{1}{(q-1)^2} \sum_{\alpha \in \mathbb{F}_q^n} S_\alpha(f)^3 \right) \\ &\leq \frac{1}{q} \left(1 + \frac{M}{(q-1)^2} \sum_{\alpha \in \mathbb{F}_q^n} S_\alpha(f)^2 \right) \\ &\leq \frac{1}{q}(1 + M) \\ &= 1 - \delta(f, \mathcal{LIN}). \end{aligned}$$

$\square$

Theorem 1.1.28 (Counterexample to the proposed cubic lower bound). *The inequality*

$$(1 - \delta(f, \mathcal{LIN}))^3 \leq \Pr[\mathbf{V}_{\text{BLR}}^f = 1]$$

does not hold for all functions $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ with the finite-field BLR verifier from Theorem 1.1.23.

Proof. It suffices to give a concrete counterexample. Take $q = 2$ and $n = 1$, so the domain and codomain are both $\mathbb{F}_2$. Let

$$f : \mathbb{F}_2 \rightarrow \mathbb{F}_2, \quad f(0) = 1, \quad f(1) = 1.$$

The homogeneous linear maps $\mathbb{F}_2 \rightarrow \mathbb{F}_2$ are exactly the zero map and the identity map. The function f agrees with the zero map at no point, and it agrees with the identity map only at the point 1. Hence

$$\delta(f, \mathcal{LIN}) = \frac{1}{2}, \quad 1 - \delta(f, \mathcal{LIN}) = \frac{1}{2}.$$

For the verifier, the only possible nonzero field elements are $a = b = 1$. Thus the verifier accepts on a pair (x, y) exactly when

$$f(x) + f(y) = f(x + y).$$

Since f is constantly 1, the left-hand side is $1 + 1 = 0$ in $\mathbb{F}_2$, while the right-hand side is 1. Therefore the verifier rejects for every choice of x, y , and

$$\Pr[V_{\text{BLR}}^f = 1] = 0.$$

But

$$(1 - \delta(f, \mathcal{LIN}))^3 = \left(\frac{1}{2}\right)^3 = \frac{1}{8} > 0 = \Pr[V_{\text{BLR}}^f = 1].$$

This contradicts the proposed lower bound. $\square$

Theorem 1.1.29 (Lower bound from agreement with a linear function). *Let $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$. Suppose there exists $\ell \in \mathcal{LIN}$ such that*

$$\Pr_{x \in \mathbb{F}_q^n} [f(x) = \ell(x)] \geq 1 - \varepsilon.$$

Then

$$\Pr[V_{\text{BLR}}^f = 1] \geq 1 - 3\varepsilon + 2\varepsilon^2.$$

In particular,

$$\Pr[V_{\text{BLR}}^f = 1] \geq 1 - 3\delta(f, \mathcal{LIN}) + 2\delta(f, \mathcal{LIN})^2.$$

Proof. Let

$$G := \{z \in \mathbb{F}_q^n : f(z) = \ell(z)\}$$

be the agreement set, and let $B := \mathbb{F}_q^n \setminus G$ be the bad set. Write

$$\eta := \Pr_z[z \in B].$$

By assumption, $\eta \leq \varepsilon$.

The verifier samples $a, b \in \mathbb{F}_q^\times$ and $x, y \in \mathbb{F}_q^n$, and queries

$$x, \quad y, \quad ax + by.$$

For fixed $a, b \in \mathbb{F}_q^\times$, the random variables x , y , and $ax + by$ are each uniform on $\mathbb{F}_q^n$. Indeed, x and y are uniform by definition, and $ax + by$ is uniform because, after fixing a , b , and x , the map

$$y \mapsto ax + by$$

is a bijection of $\mathbb{F}_q^n$.

Therefore

$$\Pr[x \in B] = \Pr[y \in B] = \Pr[ax + by \in B] = \eta.$$

Moreover the three queried points are pairwise independent. For fixed $a, b \in \mathbb{F}_q^\times$, each of the maps

$$(x, y) \mapsto (x, y), \quad (x, y) \mapsto (x, ax + by), \quad (x, y) \mapsto (y, ax + by)$$

is a bijection from $\mathbb{F}_q^n \times \mathbb{F}_q^n$ to itself. Hence

$$\Pr[x \in B, y \in B] = \Pr[x \in B, ax + by \in B] = \Pr[y \in B, ax + by \in B] = \eta^2.$$

By inclusion–exclusion and the trivial bound

$$\Pr[x \in B, y \in B, ax + by \in B] \leq \Pr[x \in B, y \in B] = \eta^2,$$

we get

$$\Pr[x \in B \text{ or } y \in B \text{ or } ax + by \in B] \leq 3\eta - 3\eta^2 + \eta^2 = 3\eta - 2\eta^2.$$

On the complementary event, the verifier sees the same three values as it would see from ℓ :

$$f(x) = \ell(x), \quad f(y) = \ell(y), \quad f(ax + by) = \ell(ax + by).$$

Since ℓ is linear,

$$\ell(ax + by) = a\ell(x) + b\ell(y).$$

Thus

$$f(ax + by) = af(x) + bf(y),$$

so the verifier accepts. Hence the acceptance probability is at least the probability of the complementary event:

$$\Pr[V_{\text{BLR}}^f = 1] \geq 1 - 3\eta + 2\eta^2.$$

If $\varepsilon \geq 1/2$, then $1 - 3\varepsilon + 2\varepsilon^2 \leq 0$, so the desired bound is trivial. Otherwise $\varepsilon \leq 1/2$, and since $\eta \leq \varepsilon$, the function $1 - 3t + 2t^2$ is decreasing on the interval containing both η and ε . This gives

$$\Pr[V_{\text{BLR}}^f = 1] \geq 1 - 3\varepsilon + 2\varepsilon^2.$$

Taking ℓ to be a closest linear function gives the final statement with $\varepsilon = \delta(f, \mathcal{LIN})$. $\square$

Theorem 1.1.30 (Boolean tightness for subspace error sets). *The bound in Theorem 1.1.29 is tight for the Boolean BLR test for infinitely many values of ε . More precisely, let $q = 2$ and let $B \leq \mathbb{F}_2^n$ be a linear subspace of density $\varepsilon = 2^{-k}$. Define*

$$f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2, \quad f(x) = \begin{cases} 1 & \text{if } x \in B, \\ 0 & \text{if } x \notin B. \end{cases}$$

Then $\delta(f, \mathcal{LIN}) = \varepsilon$ and

$$\Pr[V_{\text{BLR}}^f = 1] = 1 - 3\varepsilon + 2\varepsilon^2.$$

Proof. The zero linear function disagrees with f exactly on B , so $\delta(f, \mathcal{LIN}) \leq \varepsilon$. If $m : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$ is any nonzero linear function, then m is balanced. Thus m agrees with the indicator of B on at most half of the domain, whereas the zero linear function agrees on a $1 - \varepsilon$ fraction of the domain. For $\varepsilon \leq 1/2$, this shows that the closest linear function is the zero function, so $\delta(f, \mathcal{LIN}) = \varepsilon$.

Since $q = 2$, the verifier has $a = b = 1$ and accepts iff

$$f(x) + f(y) = f(x + y).$$

Because B is a subspace, among the three points x , y , and $x + y$, the number lying in B is never exactly two. Therefore the verifier rejects exactly when at least one of the three queried points lies in B .

The three events

$$x \in B, \quad y \in B, \quad x + y \in B$$

each have probability ε , and each pair has probability ε^2 . Since B is a subspace, the triple event also has probability ε^2 : if $x, y \in B$, then $x + y \in B$. Hence

$$\Pr[V_{\text{BLR}}^f = 0] = 3\varepsilon - 3\varepsilon^2 + \varepsilon^2 = 3\varepsilon - 2\varepsilon^2.$$

Taking complements gives

$$\Pr[V_{\text{BLR}}^f = 1] = 1 - 3\varepsilon + 2\varepsilon^2.$$

□

Chapter 2

Gowers Hatami theorem

2.1 Preliminaries

Definition 2.1.1 (σ inner product.). The σ inner product between matrices A and B is defined as

$$\langle A, B \rangle_\sigma = \text{Tr}(A^* B \sigma), \quad (2.1)$$

where $\sigma \geq 0$ and A^* denotes conjugate-transpose.

Definition 2.1.2 ((ϵ, σ) -representation.). Given a finite group G , an integer $d \geq 1, \epsilon \geq 0$, and a d -dimensional positive semidefinite matrix σ with trace 1, an (ϵ, σ) -representation of G is a map $f : G \rightarrow U_d(\mathbb{C})$, to the $d \times d$ unitary matrices, such that

$$\mathbb{E}_{x, y \in G} \Re \left(\langle f(x)^* f(y), f(x^{-1}y) \rangle_\sigma \right) \geq 1 - \epsilon, \quad (2.2)$$

where $\Re$ denotes taking real part of the inner product.

Proposition 2.1.3 (σ -norm Square). *Let σ be positive semidefinite and let A, B be matrices of the same dimension. Then*

$$\|A - B\|_\sigma^2 = \|A\|_\sigma^2 + \|B\|_\sigma^2 - 2 \Re \langle A, B \rangle_\sigma.$$

Proof. Expanding by linearity of the inner product,

$$\|A - B\|_\sigma^2 = \Re \langle A - B, A - B \rangle_\sigma = \|A\|_\sigma^2 + \|B\|_\sigma^2 - \Re \langle A, B \rangle_\sigma - \Re \langle B, A \rangle_\sigma.$$

Since σ is Hermitian, $\langle B, A \rangle_\sigma = \overline{\langle A, B \rangle_\sigma}$, so the two cross terms have the same real part, yielding $-2 \Re \langle A, B \rangle_\sigma$. $\square$

2.2 Representation theory

Definition 2.2.1 (Unitary Representation). A *unitary representation* of G is a representation to unitary matrices $\rho : G \rightarrow U(V)$.

Definition 2.2.2 (Regular representation). Let $R : G \rightarrow \mathbb{C}[G]$ be a regular representation where $\{|g\rangle, g \in G\}$ forms a basis of $\mathbb{C}[G]$. The representation is given by

$$R(x)|g\rangle = |xg\rangle. \quad (2.3)$$

It is possible to decompose any representation of a finite group into direct sums over irreps.

Theorem 2.2.3 (Maschke's theorem). *Let G be a finite group and let $\rho : G \rightarrow \text{GL}(V)$ be a finite-dimensional representation over $\mathbb{C}$. Then ρ is completely reducible. In particular, there exists a decomposition*

$$\rho \cong \bigoplus_m r_m \rho_m, \quad (2.4)$$

where $\{\rho_m\}$ is a set of inequivalent irreducible representations of G , and $r_m \in \mathbb{Z}^+$ are the multiplicities.

Proof sketch. Maschke's theorem is partially formalized in `Mathlib.RepresentationTheory.Maschke` but incomplete. Jonathan's team is formalizing this. $\square$

Proposition 2.2.4 (Decomposition of the regular representation). *Let $R : G \rightarrow \text{GL}(\mathbb{C}[G])$ be the regular representation of a finite group G . Then*

$$R \cong \bigoplus_{\rho \in \widehat{G}} d_\rho \rho, \quad (2.5)$$

where $\widehat{G}$ is a complete set of inequivalent irreducible representations of G , and $d_\rho = \dim(V_\rho)$ is the dimension of ρ .

Proof sketch. This can be proven using Maschke's theorem in 2.2.3 to decompose into irreps with multiplicity. Then use character orthogonality theorem in `mathlib.FDRep.char_orthonormal` to show the multiplicity equals irrep dimension. A more formal statement is given by the Peter-Weyl theorem. $\square$

Proposition 2.2.5 (Character of regular representation). *Let G be a finite group and $\sum_\rho$ denotes summing over the complete set of inequivalent irreps $\widehat{G}$*

$$\sum_{\rho \in \widehat{G}} d_\rho \text{Tr}(\rho(x)) = |G| \delta_{x,e}. \quad (2.6)$$

Proof sketch. The proof uses decomposition of regular representation and definition of character. $\square$

Definition 2.2.6 (Group fourier transform). Let $f : G \rightarrow U_d(\mathbb{C})$ be a map and $\rho : G \rightarrow U_{d_\rho}(\mathbb{C})$ be a unitary irrep. The fourier transform of f at the representation ρ is denoted by $\hat{f}$:

$$\hat{f}(\rho) = \mathbb{E}_{x \in G} f(x) \otimes \overline{\rho(x)}, \quad (2.7)$$

where $\overline{\rho(x)}$ denotes complex conjugate of $\rho(x)$.

2.3 Gowers Hatami Theorem

Theorem 2.3.1 (Gowers-Hatami). *Let G be a finite group, $\varepsilon \geq 0$, and $f : G \rightarrow U_d(\mathbb{C})$ an (ε, σ) -representation of G . Then there exists a $d' \geq d$, an isometry $V : \mathbb{C}^d \rightarrow \mathbb{C}^{d'}$, and a representation $g : G \rightarrow U_{d'}(\mathbb{C})$ such that*

$$\mathbb{E}_{x \in G} \|f(x) - V^* g(x) V\|_\sigma^2 \leq 2\varepsilon. \quad (2.8)$$

Proof sketch. The theorem is due to Gowers and Hatami [GH17]; the constructive proof below follows Vidick's lecture notes, Theorem 2.3 [Vid22]. In particular, the proof is constructive. The isometry is defined as $V : \mathbb{C}^d \rightarrow \mathbb{C}^d \otimes (\oplus_{\rho} \mathbb{C}^{d_{\rho}} \otimes \mathbb{C}^{d_{\rho}})$ by

$$Vu = \bigoplus_{\rho} d_{\rho}^{1/2} \sum_{i=1}^{d_{\rho}} (\hat{f}(\rho)(u \otimes e_i)) \otimes e_i. \quad (2.9)$$

And the exact representation is defined as

$$g(x) = \bigoplus_{\rho} (I_d \otimes I_{d_{\rho}} \otimes \rho(x)). \quad (2.10)$$

The proof requires the following two ingredients:

1. V is an isometry:

$$V^*V = \mathbb{I}_d \quad (2.11)$$

2. The isometry “averages” over group multiplication. For any $x \in G$,

$$V^*g(x)V = \mathbb{E}_z f(z)^* f(zx). \quad (2.12)$$

Both ingredients follow from the key identity over the regular representation in Eq. (2.6). $\square$

Next, we will give a different proof that does not directly decompose into irreps but instead uses the action of the regular representation.

Lemma 2.3.2 (Gowers-Hatami-Correlation). *Let G be a finite group, $\varepsilon \geq 0$, and $\rho : G \rightarrow U(\mathcal{H})$ an (ε, σ) -representation of G on space $\mathcal{H}$. Then there exists an isometry $V : \mathcal{H} \rightarrow \mathcal{H}'$ to a different space $\mathcal{H}'$, and a representation $\rho' : G \rightarrow \mathcal{H}'$ such that the average correlation between ρ and ρ' is:*

$$\mathbb{E}_{x \in G} \Re \langle f(x), V^* \rho'(x) V \rangle_{\sigma} \geq 1 - \varepsilon, \quad (2.13)$$

Proof. Here $\mathcal{H}' = L(G, \mathcal{H})$ is the space of functions $f : G \rightarrow \mathcal{H}$.

Let the isometry be defined by the action of the approximate representation on a state $u \in \mathcal{H}$ as

$$V : \mathcal{H} \rightarrow L(G, \mathcal{H}), \quad (2.14)$$

$$V(u) : G \rightarrow \mathcal{H}, \quad x \mapsto \rho(x)(u), \forall x \in G. \quad (2.15)$$

The adjoint map V^* is defined as the unique map such that for all $u \in \mathcal{H}$ and $g \in L(G, \mathcal{H})$,

$$\langle V(u), g \rangle_{L(G, \mathcal{H})} = \langle u, V^*(g) \rangle_{\mathcal{H}}. \quad (2.16)$$

In particular, V^* is explicitly given by the group average

$$V^* : L(G, \mathcal{H}) \rightarrow \mathcal{H}, \quad (2.17)$$

$$V^*(f) = \mathbb{E}_{x \in G} [\rho^*(x) f(x)]. \quad (2.18)$$

The exact representation is given by the right regular representation on $L(G, \mathcal{H})$:

$$\rho' : G \rightarrow \text{End}(L(G, \mathcal{H})), \quad (2.19)$$

$$(\rho'(x)f)(y) = f(yx), \quad \forall x, y \in G. \quad (2.20)$$

V is an isometry because

$$V^*V(u) = \mathbb{E}_{x \in G} [\rho^*(x)\rho(x)(u)] = u. \quad (2.21)$$

The isometry “averages” over group multiplication because

$$\begin{aligned} V^*\rho'(x)V(u) &= \mathbb{E}_{y \in G} [\rho^*(y)(\rho'(x)\rho(y))(u)] \\ &= \mathbb{E}_{y \in G} [\rho^*(y)(\rho(yx)(u))]. \end{aligned} \quad (2.22)$$

Finally, using the isometry average property in Eq. 2.22 and the definition of (ε, σ) -representation, we have

$$\begin{aligned} &\mathbb{E}_{x \in G} \Re \langle \rho(x), V^*\rho'(x)V \rangle_\sigma, \\ &= \mathbb{E}_{x, y \in G} \Re \langle \rho(x), \rho^*(y)\rho(yx) \rangle_\sigma, \\ &\geq 1 - \varepsilon. \end{aligned} \quad (2.23)$$

□

Theorem 2.3.3 (Gowers-Hatami-2). *Let G be a finite group, $\varepsilon \geq 0$, and $\rho : G \rightarrow U(\mathcal{H})$ an (ε, σ) -representation of G on space $\mathcal{H} = \mathbb{C}^d$. Then there exists an isometry $V : \mathcal{H} \rightarrow \mathcal{H}'$ to a different space $\mathcal{H}'$, and a representation $\rho' : G \rightarrow \mathcal{H}'$ such that*

$$\mathbb{E}_{x \in G} \|\rho(x) - V^*\rho'(x)V\|_\sigma^2 \leq 2\varepsilon. \quad (2.24)$$

Proof. A regular-representation proof of the Gowers-Hatami theorem appears in Morel’s slides [Mor20], following the original Gowers-Hatami theorem [GH17].

The rest of the proof follows by expanding the σ -norm, bounding the norm of $\mathbb{E}_{x \in G} V^*\rho'(x)V$ and applying the definition of (ε, σ) -representation. In particular, we have

$$\begin{aligned} &\mathbb{E}_{x \in G} \|\rho(x) - V^*\rho'(x)V\|_\sigma^2 \\ &= \mathbb{E}_{x \in G} \|\rho(x)\|_\sigma^2 + \mathbb{E}_{x \in G} \|V^*\rho'(x)V\|_\sigma^2 - 2\mathbb{E}_{x \in G} \Re \langle \rho(x), V^*\rho'(x)V \rangle_\sigma \end{aligned} \quad (2.25)$$

Because ρ is unitary and $\text{Tr} \sigma = 1$, $\mathbb{E}_{x \in G} \|\rho(x)\|_\sigma^2 = 1$. For the second term, recall from Eq. (2.22) that the compression is the group average $V^*\rho'(x)V = \mathbb{E}_{y \in G} \rho^*(y)\rho(yx)$. Each summand $\rho^*(y)\rho(yx)$ is a product of unitaries, hence unitary, so $\|\rho^*(y)\rho(yx)\|_\sigma = 1$. By convexity of the σ -norm—the triangle inequality applied to the average—we obtain

$$\|V^*\rho'(x)V\|_\sigma = \|\mathbb{E}_{y \in G} \rho^*(y)\rho(yx)\|_\sigma \leq \mathbb{E}_{y \in G} \|\rho^*(y)\rho(yx)\|_\sigma = 1, \quad (2.26)$$

hence $\mathbb{E}_{x \in G} \|V^*\rho'(x)V\|_\sigma^2 \leq 1$.

Therefore using the bounds in Eq. 2.26 and Eq. 2.23 by applying Lemma 2.3.2, we have

$$\mathbb{E}_{x \in G} \|\rho(x) - V^*\rho'(x)V\|_\sigma^2 \leq 2 - 2(1 - \varepsilon) = 2\varepsilon. \quad (2.27)$$

□

2.4 Gowers-Hatami for BLR linearity test over $\mathbb{F}_2$

In classical BLR, the mapping is from a field $\mathbb{F}_2^n$. First let us relate the BLR test to the notion of approximate representation.

Proposition 2.4.1 (BLR test induces an approximate representation). *Let $f : \mathbb{Z}_2^n \rightarrow \mathbb{Z}_2$ satisfy*

$$\Pr_{x,y}[f(x) + f(y) = f(x+y)] \geq 1 - \epsilon.$$

Define $F : \mathbb{Z}_2^n \rightarrow \{\pm 1\}$ by $F(x) = (-1)^{f(x)}$. Then

$$\mathbb{E}_{x,y}|F(x)F(y) - F(x+y)|^2 \leq 4\epsilon,$$

so F is an $(2\epsilon, 1)$ -approximate representation.

Proof. First observe that for all $x, y \in G$,

$$F(x)F(y) = (-1)^{f(x)+f(y)}.$$

Hence

$$F(x)F(y) = F(x+y) \iff f(x) + f(y) = f(x+y) \pmod{2}.$$

Define the indicator of BLR success:

$$\mathbf{1}_{\text{good}}(x, y) := \mathbf{1}[f(x) + f(y) = f(x+y)].$$

Then

$$\Pr_{x,y}[\mathbf{1}_{\text{good}} = 1] \geq 1 - \epsilon.$$

Now compute the squared error. Since $F(x) \in \{\pm 1\}$,

$$|F(x)F(y) - F(x+y)|^2 = \begin{cases} 0 & \text{if } F(x)F(y) = F(x+y), \\ 4 & \text{otherwise.} \end{cases}$$

Taking expectation over (x, y) ,

$$\begin{aligned} \mathbb{E}_{x,y}|F(x)F(y) - F(x+y)|^2 &= 4 \Pr_{x,y}[F(x)F(y) \neq F(x+y)] \\ &= 4 \Pr_{x,y}[f(x) + f(y) \neq f(x+y)] \\ &\leq 4\epsilon. \end{aligned}$$

□

The Gowers-Hatami theorem can be applied to show that if a function $f : \mathbb{Z}_2^n \rightarrow \mathbb{F}_2$ passes the BLR linearity test with high probability, then there exists a linear function $\ell : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$ that is close to f in terms of Hamming distance $\delta(f, \ell)$ defined in Def. 1.1.15. To do this, we first need to apply the Gowers-Hatami theorem to the approximate representation $F(x) = (-1)^{f(x)}$ given by Proposition 2.4.1.

Corollary 2.4.2 (Gowers-Hatami Correlation for BLR). *Let $F : \mathbb{Z}_2^n \rightarrow U(1)$ be an (ϵ, σ) approximate representation. Then there exists a unit vector v and an exact representation G such that*

$$\mathbb{E}_{x \in \mathbb{Z}_2^n} \Re \langle F(x), \langle v, G(x)v \rangle \rangle \geq 1 - \epsilon. \quad (2.28)$$

Next, we would like to conclude from the above corollary that there exists a linear function $\ell : \mathbb{Z}_2^n \rightarrow \mathbb{Z}_2$ such that $\Pr_x[f(x) \neq \ell(x)] \leq \epsilon/2$. To do this, we need to use a fact that about the regular representation of $\mathbb{Z}_2^n$.

Fact 2.4.3 (Regular representation of $\mathbb{Z}_2^n$). *Let $\rho(a)$ act on $L = \{f : \mathbb{Z}_2^n \rightarrow \mathbb{C}\}$ by $(\rho(a)f)(x) = f(x + a)$. Then in the Fourier basis $\{|S\rangle\}$, we have*

$$\rho(a) = \sum_{S \subseteq [n]} \chi_S(a) |S\rangle \langle S|.$$

Lemma 2.4.4 (Gowers-Hatami implies near linearity). *Let $F : \mathbb{Z}_2^n \rightarrow \mathbb{C}$ be an $(\epsilon, 1)$ approximate representation of $\mathbb{Z}_2^n$ from the BLR test such that $F(x) = (-1)^{f(x)}$ for some function $f : \mathbb{Z}_2^n \rightarrow \mathbb{Z}_2$. Then there exists a linear function $\ell : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$ such that the distance between f and ℓ is bounded*

$$\delta(f, \ell) \leq \epsilon/2. \quad (2.29)$$

Proof. By Corollary 2.4.2, there exists a representation G and a unit vector v such that

$$\mathbb{E}_x \Re \langle F(x), \langle v, G(x)v \rangle \rangle \geq 1 - \epsilon.$$

Using the diagonal form of $G(x)$ in the character basis,

$$G(x) = \sum_S \chi_S(x) |S\rangle \langle S|.$$

The inner product condition can be rewritten as

$$\sum_S p_S \mathbb{E}_x \Re \langle F(x), \chi_S(x) \rangle \geq 1 - \epsilon,$$

where $p_S := |\langle v | S \rangle|^2$ and $\sum_S p_S = 1$. Therefore,

$$\sum_S p_S \widehat{F}(S) \geq 1 - \epsilon,$$

where $\widehat{F}(S) := \mathbb{E}_x [F(x) \chi_S(x)]$ is the real-valued Fourier coefficient of F at S .

Since $(p_S)_S$ is a probability distribution, there exists some S^* such that

$$\widehat{F}(S^*) \geq 1 - \epsilon.$$

This means that there is a large Fourier coefficient at S^* . Define $\ell(x) = x \cdot S^*$. Then $\chi_{S^*}(x) = (-1)^{\ell(x)}$, and

$$\mathbb{E}_x [F(x) \chi_{S^*}(x)] = 1 - 2 \Pr_x[f(x) \neq \ell(x)].$$

Hence

$$1 - 2\delta(f, \ell) \geq 1 - \epsilon,$$

which implies

$$\delta(f, \ell) \leq \epsilon/2.$$

□

Proposition 2.4.5 (Equivalence for BLR test). *Let $\Pr[V_{BLR}^f = 0]$ denote the probability that the BLR verifier rejects a function $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$. Let $\delta(f, \mathcal{LIN})$ denote the distance between f and the closest linear function in $\mathcal{LIN}$. Then the following statements are equivalent:*

- If $\Pr[V_{BLR}^f = 0] \leq \epsilon$, then $\delta(f, \mathcal{LIN}) \leq \epsilon$.
- $\Pr[V_{BLR}^f = 0] \geq \delta(f, \mathcal{LIN})$.

Now we provide the proof of Soundness of BLR test using the Gowers-Hatami theorem.

Theorem 2.4.6 (Soundness of the Z_2 BLR). *For every function $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$,*

$$\Pr[V_{BLR}^f = 0] \geq \delta(f, \mathcal{LIN}).$$

Equivalently,

$$\Pr[V_{BLR}^f = 1] \leq 1 - \delta(f, \mathcal{LIN}).$$

Proof. Let $F : \mathbb{Z}_2^n \rightarrow \{-1, 1\}$ be defined by $F(x) = (-1)^{f(x)}$, as given by Proposition 2.4.1. F is an $(2\epsilon, 1)$ approximate representation of $\mathbb{Z}_2^n$. Applying Gowers-Hatami theorem by Lemma 2.4.4, then there exists a linear function ℓ such that $\delta(f, \ell) \leq \epsilon$ so that $\delta(f, \mathcal{LIN}) \leq \epsilon$.

This proves the implication $\Pr[V_{BLR}^f = 0] \leq \epsilon \Rightarrow \delta(f, \mathcal{LIN}) \leq \epsilon$, which is equivalent to

$$\Pr[V_{BLR}^f = 0] \geq \delta(f, \mathcal{LIN}),$$

by Proposition 2.4.5. □

2.5 Gowers-Hatami for BLR linearity test over $\mathbb{F}_q$

Assume $q = p^s$ be a prime with power $s \in \mathbb{Z}^+$, $n \geq 1$, and we set

$$X := \mathbb{F}_q^n, \quad A := \mathbb{F}_q^\times, \quad \omega := e^{2\pi i/p}.$$

Definition 2.5.1 (BLR verifier and distance). Given $f : X \rightarrow \mathbb{F}_q$, the verifier V_{BLR}^f samples $a, b \in A$ and $x, y \in X$ uniformly and accepts iff $af(x) + bf(y) = f(ax + by)$. Let the test rejection probability be bounded by

$$\Pr_{a,b,x,y} [af(x) + bf(y) \neq f(ax + by)] \leq \epsilon.$$

For $\alpha \in X$, write $\ell_\alpha(x) := \langle \alpha, x \rangle$, $\delta_\alpha := \Pr_x[f(x) \neq \ell_\alpha(x)]$, and $\delta(f, \mathcal{LIN}) := \min_\alpha \delta_\alpha$.

In the rest of the section, we prove the BLR soundness using Gowers-Hatami theorem.

Theorem 2.5.2 (Soundness of the BLR test). *For every $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, $\delta(f, \mathcal{LIN}) \leq \Pr_{a,b,x,y} [af(x) + bf(y) \neq f(ax + by)]$.*

First notice that linearity implies two properties:

1. Additivity $f(x + y) = f(x) + f(y)$, $\forall x, y \in X$.

Note that for q prime additivity is actually equivalent to linearity. However this is not the case for prime powers $q = p^s$ for $s > 1$.

2. Homogeneity $f(ax) = af(x)$, $\forall x \in X, a \in A$.

For $q = 2$, the two properties are redundant. For $q > 2$, turns out testing homogeneity is crucial for a test to have a high soundness.

In the rest of the section, we discuss an encoding of both properties through a group construction by ChatGPT.

2.5.1 The twisted group G

Define a group G by the set $X \times A$ with the binary operation

$$(x, a) \star (y, b) := (b^{-1}x + a^{-1}y, ab). \quad (2.30)$$

Lemma 2.5.3 (G is an abelian group). *The operation $\star$ in (2.30) makes G into an abelian group with identity $e = (0, 1)$ and inverses $(x, a)^{-1} = (-a^2x, a^{-1})$.*

Proof. Associativity. For $(x, a), (y, b), (z, c) \in G$,

$$((x, a) \star (y, b)) \star (z, c) = (b^{-1}x + a^{-1}y, ab) \star (z, c) = (b^{-1}c^{-1}x + a^{-1}c^{-1}y + a^{-1}b^{-1}z, abc),$$

and an identical computation gives the same expression for $(x, a) \star ((y, b) \star (z, c))$.

Commutativity. $(y, b) \star (x, a) = (a^{-1}y + b^{-1}x, ba) = (b^{-1}x + a^{-1}y, ab) = (x, a) \star (y, b)$.

Identity. For any $(y, b) \in G$, $(0, 1) \star (y, b) = (b^{-1} \cdot 0 + 1^{-1} \cdot y, 1 \cdot b) = (y, b)$.

Inverse. For any $(x, a) \in G$,

$$(x, a) \star (-a^2x, a^{-1}) = ((a^{-1})^{-1} \cdot x + a^{-1} \cdot (-a^2x), a \cdot a^{-1}) = (ax - ax, 1) = (0, 1).$$

□

It is at first surprising to come up with $\star$ operation, but G is isomorphic to the direct product group.

Proposition 2.5.4 (Group isomorphism to product group). *Let $(G, \star)$ be an abelian group defined by Equation 2.30. The map*

$$\Psi : G \rightarrow (X, +) \times (A, \cdot), \quad \Psi(x, a) := (ax, a), \quad (2.31)$$

is a group isomorphism, where the codomain is the direct product of $(X, +)$ and $(A, \cdot)$ with operation $\star_\Psi : (x, a) \star_\Psi (y, b) = (x + y, ab)$.

Proof. Compute

$$\Psi((x, a) \star (y, b)) = \Psi(b^{-1}x + a^{-1}y, ab) = (ab \cdot (b^{-1}x + a^{-1}y), ab) = (ax + by, ab),$$

which equals $\Psi(x, a) \star_\Psi \Psi(y, b) = (ax + by, ab)$. The map Ψ is bijective with inverse $\Psi^{-1}(y, b) = (b^{-1}y, b)$. □

The isomorphism gives us a easy way to compute its irreducible representations which forms the Pontryagin dual group (the set of irrep characters for an abelian group).

Fact 2.5.5 (Pontryagin dual over direct product). *Let G_1, G_2 be finite abelian groups. Let $H = G_1 \times G_2$ be their direct product. Then the Pontryagin dual formed by all the irreps. is $\hat{H} = \hat{G}_1 \times \hat{G}_2$.*

Proposition 2.5.6 (Group characters). *The group $(G, \star)$ has the Pontryagin dual character group*

$$\hat{G} = \{\chi_{\alpha, k}\}_{\alpha \in X, k \in [0, \dots, q-2]}, \quad \chi_{\alpha, k}(x, a) = \omega^{\text{Tr}(\alpha, ax)} \psi_k(a), \quad (2.32)$$

where $\psi_k(a) = e^{\frac{i2\pi mk}{q-1}}$ and $a = g_0^m$ for a fixed multiplicative generator $g_0 \in A$.

Proof. By Proposition 2.5.4, $(G, \star) \cong (X, +) \times (A, \cdot)$ via $\Psi(x, a) = (ax, a)$. First we consider the characters of the product group. Using Fact. 2.5.5, we have

$$(X, +) \times (A, \cdot) = \hat{X} \times \hat{A}.$$

The character groups of X and A are

$$\hat{X} = \{\phi_\alpha(y) = \omega^{\text{Tr}(\alpha, y)}\}_{\alpha \in X}, \quad (2.33)$$

$$\hat{A} = \{\psi_k(b) = e^{\frac{i2\pi mk}{q-1}}\}_{k \in [0, \dots, q-2]}, \quad (2.34)$$

where $b = g_0^m$ for a fixed multiplicative generator $g_0 \in A$. Here the field trace $\text{Tr} : \mathbb{F}_q \rightarrow \mathbb{F}_p$ is defined by $\text{Tr}(x) = x + x^p + x^{p^2} + \dots + x^{p^{s-1}}$ for $x \in \mathbb{F}_q$ (see Definition 1.1 in Chapter 1). Therefore, $(X, +) \times (A, \cdot) = \{\Theta_{\alpha, k}(y, b) = \phi_\alpha(y)\psi_k(b)\}_{\alpha, k}$. The character $\chi \in \hat{G}$ is given by the pull-back through the isomorphism

$$\chi_{\alpha, k}(x, a) = \Theta_{\alpha, k}(\Psi(x, a)) = \Theta_{\alpha, k}(ax, a). \quad (2.35)$$

Therefore, we have $\chi_{\alpha, k}(x, a) = \phi_\alpha(ax)\psi_k(a) = \omega^{\text{Tr}(\alpha, ax)}\psi_k(a)$. $\square$

2.5.2 Approximate representation and matrix lift of f

Define a unitary matrix valued lift $F_f : G \rightarrow U(\mathbb{C}^{q-1})$ by the diagonal matrix

$$F_f(x, a) = D(af(x)) := \sum_{c \in A} \omega^{\text{Tr } caf(x)} |c\rangle\langle c|, \quad (2.36)$$

where $\{|c\rangle\}$ denotes the orthonormal basis vectors of a Hilbert space $\mathcal{H} \cong \mathbb{C}^{q-1}$ and the inner product between operators defined by $\langle A, B \rangle_\sigma = \text{Tr}(A^* B \sigma)$ and choose $\sigma = \mathbb{I}/(q-1)$. We first show that the BLR condition indeed gives an approximate representation.

Lemma 2.5.7 (Lift operation from BLR condition). *Let $f : X \rightarrow \mathbb{F}_q$ and $(x, a), (y, b) \in G$ be rejected by the BLR test with probability ε , then the unitary matrix lift F_f in Equation 2.36 is an $(\frac{q}{q-1}\varepsilon, \frac{\mathbb{I}}{q-1})$ approximate representation of $(G, \star)$.*

Proof. The BLR test condition gives

$$\Pr_{c, d, x, y} [cf(x) + df(y) \neq f(cx + dy)] \leq \varepsilon.$$

First we observe that

$$\Pr_{g, h \in G} [F_f(g \star h) = F_f(g)F_f(h)] = \Pr_{c, d, x, y} [cf(x) + df(y) = f(cx + dy)]. \quad (2.37)$$

Using Equation 2.30 and Equation 2.36,

$$\begin{aligned} F_f((x, a) \star (y, b)) &= \sum_{c \in A} \omega^{\text{Tr } cab f(b^{-1}x + a^{-1}y)} |c\rangle\langle c|, \\ F_f(x, a) F_f(y, b) &= \sum_{c \in A} \omega^{\text{Tr } c(af(x) + bf(y))} |c\rangle\langle c|. \end{aligned}$$

These coincide iff their exponents agree modulo p for $c \in A$:

$$ab f(b^{-1}x + a^{-1}y) = af(x) + bf(y) \pmod{p}, \quad \forall c \in A.$$

Therefore, we have equality over $\mathbb{F}_p$

$$abf(b^{-1}x + a^{-1}y) = af(x) + bf(y).$$

Since a, b have a multiplicative inverse, substitute $c := b^{-1}$, $d := a^{-1}$, we have

$$f(cx + dy) = cf(x) + df(y),$$

which proves Equation 2.37.

To relate the test condition to the approximate representation, first note the equivalent definition of approximate representation

$$\mathbb{E}_{g,h \in G} \Re \langle F_f(g)^* F_f(h), F_f(g^{-1}h) \rangle_\sigma = \mathbb{E}_{g,h \in G} \Re \langle F_f(g) F_f(h), F_f(gh) \rangle_\sigma, \quad (2.38)$$

because F_f is unitary so $F_f(g)^* = F_f(g^{-1})$, $\forall g \in G$. Then evaluate the right hand side

$$\begin{aligned} \mathbb{E}_{g,h \in G} \Re \langle F_f(g) F_f(h), F_f(gh) \rangle_\sigma &= \mathbb{E}_{g,h \in G} \Re \text{Tr} [F_f^*(h) F_f^*(g) F_f(gh) \sigma] , \\ &= \mathbb{E}_{(x,a),(y,b) \in G} \frac{1}{q-1} \Re \text{Tr} \left[\sum_{c \in A} \omega^{-\text{Tr } c(af(x)+bf(y))} \omega^{\text{Tr } cabf(b^{-1}x+a^{-1}y)} |c\rangle \langle c| \right] , \\ &= \mathbb{E}_{(x,a),(y,b) \in G} \frac{1}{q-1} \Re \sum_{c \in A} \omega^{\text{Tr } cz}, \end{aligned} \quad (2.39)$$

where we define $z = c(abf(b^{-1}x + a^{-1}y) - (af(x) + bf(y)))$. To compute the expectation, we can split to two cases

$$\sum_{c \in A} \omega^{\text{Tr } cz} = \begin{cases} q-1, & z = 0 \quad \text{when } f \text{ passes BLR,} \\ -1, & z \neq 0 \quad \text{when } f \text{ fails BLR.} \end{cases} \quad (2.40)$$

The second case is because $\sum_{c \in A} \omega^{\text{Tr } cz} + 1 = \sum_{c \in \mathbb{F}_q} \omega^{\text{Tr } cz} = 0$. Therefore, we have the approximate representation

$$\mathbb{E}_{g,h \in G} \Re \langle F_f(g) F_f(h), F_f(gh) \rangle_\sigma \geq (1 - \varepsilon) - \frac{1}{q-1} \varepsilon = 1 - \frac{q}{q-1} \varepsilon. \quad (2.41)$$

Hence F_f is an approximate representation with $\epsilon = \frac{q}{q-1} \varepsilon$ as claimed. $\square$

The approximate representation “appears” slightly worse than what we would have hoped $\epsilon > \varepsilon$, but we will see this is not an issue for achieving a high soundness. Remarkably, the twisted group operation let us recover the linearity test condition exactly.

Fact 2.5.8. *F_f is an exact representation of $(G, \star)$ if and only if $f \in \mathcal{LIN}$.*

Proof. If F_f is an exact representation of G , then Equation 2.37 forces $\varepsilon = 0$. Setting $c = d = 1$ in the BLR condition gives $f(x + y) = f(x) + f(y)$ for all x, y , so f is additive. Set $y = 0$ and $c = 1 \in A$: $f(x) = f(x) + d f(0)$ for any $d \in A$. Set $y = 0$ and $d = 1$, then $f(0) = 0$ and $f(cx) = cf(x)$. Combined with additivity, f is $\mathbb{F}_q$ -linear, so $f \in \mathcal{LIN}$.

Conversely, if $f = \ell_\alpha$ for some $\alpha \in X$, then so F_f is a group homomorphism and an exact representation of G . $\square$

2.5.3 Gowers-Hatami implies BLR high soundness

In this section, we first prove the key lemma relating the implication of GH theorem to the hamming distance. We provide two proofs:

1. The first proof in Lemma 2.5.9 should only use the existence of an isometry V and exact representation R . Unfortunately, we ended up with an extra constraint about the range of ϵ .
2. In the second proof in Lemma 2.5.11, we use an explicit isometry V constructed via F_f in the proof of Gowers-Hatami theorem. This proof might require a extra care in formalization of lean. However, this hybrid Fourier-GH approach gives us an unconditional high soundness of the linear BLR test.

Finally, we provide a proof of Theorem 2.5.1.

Lemma 2.5.9 (Gowers-Hatami to linearity (Convex character proof)). *Suppose $F_f : G \rightarrow U(\mathbb{C}^{q-1})$ defined in Equation 2.36 via $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ is an (ϵ, σ) approximate unitary representation of $(G, \star)$ defined in Equation 2.30. Then, there exists a linear function $\ell : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ such that the relative Hamming distance $\delta(f, \ell)$ is bounded by:*

$$\delta(f, \ell) \leq \frac{q-1}{q}\epsilon. \quad (2.42)$$

Proof. By Gowers-Hatami theorem, we have an isometry $V : \mathcal{H} \rightarrow \mathcal{H}'$ and an exact representation $R : G \rightarrow \text{End}(\mathcal{H}')$, where the Hilbert space $\mathcal{H} \cong \mathbb{C}^{q-1}$ and the space R acts on $\mathcal{H}' = L(G, \mathcal{H})$

$$\mathbb{E}_{g \in G} \Re \langle F_f(g), V^* R(g) V \rangle_\sigma \geq 1 - \epsilon. \quad (2.43)$$

Let $\{|c\rangle\}_{c \in \mathbb{F}_q^\times}$ be the orthonormal basis for $\mathcal{H}$, and the set of characters $\{\chi\}_{\chi \in \hat{G}}$ be an orthonormal basis for $L(G, \mathbb{C})$.

Because $F_f(g)$ is diagonal in $\{|c\rangle\}$ with entries $\omega^{\text{Tr } c a f(x)}$, we can rewrite the inner product trace:

$$\langle F_f(g), V^* R(g) V \rangle_\sigma = \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \langle c | F_f^*(g) V^* R(g) V | c \rangle = \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \omega^{-\text{Tr } c a f(x)} \langle v_c, R(g) v_c \rangle_{L(G, \mathcal{H})},$$

where $|v_c\rangle = V|c\rangle \in \mathcal{H}'$. Taking the expectation over $g = (x, a)$ and taking the real part, condition Equation 2.62 becomes:

$$\frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x, a} \Re [\omega^{-\text{Tr } c a f(x)} \langle v_c, R(x, a) v_c \rangle] \geq 1 - \epsilon. \quad (2.44)$$

A basis for $L(G, \mathcal{H})$ is given by $\{\chi|c\rangle\}_{\chi \in \hat{G}, c \in \mathbb{F}_q^\times}$. By Proposition 2.5.6, the characters are given by

$$\hat{G} = \{\chi_{\alpha, k}\}_{\alpha \in X, k \in [0, \dots, q-2]}, \chi_{\alpha, k}(x, a) = \omega^{\text{Tr}(\alpha, ax)} \psi_k(a).$$

The regular representation $R(g)$ acts on $L(G, \mathcal{H})$ so it can be decomposed as

$$R(g) = \sum_{\alpha, k} \chi_{\alpha, k}(g) P_{\alpha, k}, \quad (2.45)$$

where $P_{\alpha, k} = |\chi_{\alpha, k}\rangle \langle \chi_{\alpha, k}| \otimes \mathbb{I}_{\mathcal{H}}$ projects onto the character function subspace (this leaves $\mathcal{H}$ invariant).

Substituting this into the matrix element for v_c , we obtain:

$$\langle v_c, R(g)v_c \rangle = \sum_{\alpha, k} \chi_{\alpha, k}(g) \langle v_c, P_{\alpha, k} v_c \rangle.$$

Let $p_{c, \alpha, k} := \langle v_c, P_{\alpha, k} v_c \rangle$. As $P_{\alpha, k}$ is a projector, then

$$p_{c, \alpha, k} = \langle v_c | \chi_{\alpha, k} \rangle \langle \chi_{\alpha, k} | \otimes \mathbb{I}_{\mathcal{H}} | v_c \rangle \geq 0. \quad (2.46)$$

Since V is an isometry $V^*V = \mathbb{I}_{\mathcal{H}}$, we have $\langle c | V^*V | c \rangle = \langle c, c \rangle_{\mathcal{H}} = 1$. Equivalently, we have $\langle v_c, v_c \rangle_{L(G, \mathcal{H})} = 1$. Since $\sum P_{\alpha, k} = I_{L(G, H)}$, it follows that $p_{c, \alpha, k}$ is a valid probability distribution over the characters,

$$\sum_{\alpha, k} p_{c, \alpha, k} = 1, \quad \forall c \in \mathbb{F}_q^\times. \quad (2.47)$$

Substituting the spectral decomposition Equation 2.65 back into Equation 2.63, we get:

$$\sum_{c \in \mathbb{F}_q^\times} \left(\frac{1}{q-1} \right) \sum_{\alpha, k} p_{c, \alpha, k} \Re \left(\mathbb{E}_{x, a} \left[\omega^{-\text{Tr } c a f(x)} \chi_{\alpha, k}(x, a) \right] \right) \geq 1 - \epsilon.$$

This is a convex combination (a weighted average) of real values less than 1. Because the average is at least $1 - \epsilon$, there must exist at least one specific configuration (c^*, α^*, k^*) with $c^* \neq 0$ such that its individual value satisfies:

$$\Re \left(\mathbb{E}_{x, a} \left[\omega^{-\text{Tr } c^* a f(x)} \omega^{\text{Tr } \langle \alpha^*, a x \rangle} \psi_{k^*}(a) \right] \right) \geq 1 - \epsilon. \quad (2.48)$$

Because $c^* \in \mathbb{F}_q^\times$ is invertible, we define a specific linear function $\ell(x) := \langle (c^*)^{-1} \alpha^*, x \rangle$. Notice that $\langle \alpha^*, a x \rangle = c^* a \ell(x)$. We can rewrite the argument of the expectation in Equation 2.72 as:

$$E^* := \mathbb{E}_x \left[\mathbb{E}_a \left[\omega^{\text{Tr } c^* a (\ell(x) - f(x))} \psi_{k^*}(a) \right] \right].$$

For the rest of this proof, without using specific form of $p_{c, \alpha, k}$, we consider two cases: $k^* = 0$ and $k^* \neq 0$. Define $z(x) = c^*(\ell(x) - f(x))$.

1. Consider $k^* = 0$ so that $\psi_0(a) = 1$, the inner expectation over $a \in \mathbb{F}_q^\times$ for a fixed x :

$$\mathbb{E}_{a \in \mathbb{F}_q^\times} \omega^{\text{Tr } c^* a (\ell(x) - f(x))} = \begin{cases} 1, & z(x) = 0, \\ -\frac{1}{q-1}, & z(x) \neq 0, \end{cases} \quad (2.49)$$

where in the second case we have again used average of roots of unity is $\sum_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b} + 1 = 0$.

Evaluating the outer expectation over x and recall the definition of hamming distance

$$\delta_\alpha := \Pr_x[f(x) \neq \ell_\alpha(x)], \quad (2.50)$$

we find:

$$\Re E^* = E^* = 1 \cdot (1 - \delta(f, \ell)) + \left(\frac{-1}{q-1} \right) \cdot \delta(f, \ell) = 1 - \delta(f, \ell) \left(1 + \frac{1}{q-1} \right) = 1 - \delta(f, \ell) \left(\frac{q}{q-1} \right).$$

Applying our bound from Equation 2.72, we have:

$$1 - \delta(f, \ell) \left(\frac{q}{q-1} \right) \geq 1 - \epsilon.$$

Rearranging yields the final result:

$$\delta(f, \ell) \leq \frac{q-1}{q} \epsilon.$$

2. Consider $k^* \neq 0$, then by Proposition 2.5.6 we have $\psi_k(a) = e^{\frac{i2\pi mk}{q-1}}$ and $a = g_0^m$ for a fixed multiplicative generator $g_0 \in \mathbb{F}_q^\times$. Note that averaging over $\psi_k(a)$ we have

$$\mathbb{E}_{a \in \mathbb{F}_q^\times} \psi_k(a) = \frac{1}{q-1} \sum_{m=0}^{q-2} e^{\frac{i2\pi mk}{q-1}} = 0. \quad (2.51)$$

The average expectation over $a \in \mathbb{F}_q^\times$ for a fixed x gives:

$$\mathbb{E}_{a \in \mathbb{F}_q^\times} \omega^{\text{Tr } c^* a(\ell(x) - f(x))} \psi_{k^*}(a) = \begin{cases} 0, & z(x) = 0, \\ \frac{1}{q-1} \sum_{a \in \mathbb{F}_q^\times} \omega^{\text{Tr } az(x)} \psi_{k^*}(a), & z(x) \neq 0. \end{cases} \quad (2.52)$$

In the second case, let $b = az(x)$ so that

$$\begin{aligned} \frac{1}{q-1} \sum_{a \in \mathbb{F}_q^\times} \omega^{\text{Tr } az(x)} \psi_{k^*}(a) &= \frac{1}{q-1} \sum_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b} \psi_{k^*}(z^{-1}b), \\ &= \frac{\psi_{k^*}(z^{-1})}{q-1} \sum_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b} \psi_{k^*}(b). \end{aligned} \quad (2.53)$$

It turns out the summation is a known quantity called *the Gauss sum*

$$\mathcal{G}(\psi_{k^*}) = \sum_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b} \psi_{k^*}(b). \quad (2.54)$$

Is it believed that there is no analytic expression for $\mathcal{G}$ but it is known that

$$|\mathcal{G}(\psi_{k^*})| = \sqrt{q} \quad (2.55)$$

Evaluating the outer expectation over x and recall the definition of hamming distance gives

$$\Re E^* = 0 \cdot (1 - \delta(f, \ell)) + \Re \left(\frac{\psi_{k^*}(z^{-1})}{q-1} \mathcal{G}(\psi_{k^*}) \right) \cdot \delta(f, \ell).$$

Therefore we have

$$\Re E^* = \Re \left(\frac{\psi_{k^*}(z^{-1})}{q-1} \mathcal{G}(\psi_{k^*}) \right) \delta(f, \ell) \leq \frac{|\mathcal{G}(\psi_{k^*})|}{q-1} = \frac{\sqrt{q}}{q-1}.$$

On the other hand, from Equation 2.72 we have

$$\Re E^* \geq 1 - \epsilon.$$

This is a contradiction to the existence of k^* if

$$\frac{\sqrt{q}}{q-1} < 1 - \epsilon. \quad (2.56)$$

Then if we have a good test outcome such that Equation 2.56 is satisfied, then we only have a contribution from $k^* = 0$, where we have already achieved the bound.

□

The above bound is tight assuming we pass the test with high probability such that $\epsilon < 1 - \frac{\sqrt{q}}{q-1}$. Next, we try to give a proof that is unconditional, using Fourier coefficients.

First let us write down statements about Fourier coefficients of F_f , which we will use in the proof.

Proposition 2.5.10 (Fourier coefficient of F_f). *Let the matrix lift $F_f : G \rightarrow U(\mathbb{C}^{q-1})$ be defined by the diagonal matrix*

$$F_f(x, a) = D(af(x)) := \sum_{c \in A} \omega^{\text{Tr } caf(x)} |c\rangle \langle c|,$$

then its Fourier coefficients can be written as

$$[\hat{F}_f(\chi_{c\beta, k})]_{cc} = \psi_k(c) \nu_k(\beta), \quad (2.57)$$

where

$$\nu_k(\beta) := \mathbb{E}_{x, b} \left[\omega^{\text{Tr } b(f(x) - \langle \beta, x \rangle)} \overline{\psi_k(b)} \right], \quad (2.58)$$

Proof. By Proposition 2.5.6, the characters are given by

$$\hat{G} = \{\chi_{\alpha, k}\}_{\alpha \in X, k \in [0, \dots, q-2]}, \chi_{\alpha, k}(x, a) = \omega^{\text{Tr}(\alpha, ax)} \psi_k(a).$$

Therefore the Fourier coefficients of F_f are

$$\begin{aligned} \hat{F}_f(\chi_{\alpha, k}) &= \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x, a} \omega^{\text{Tr } caf(x)} \bar{\chi}_{\alpha, k}(x, a), \\ &= \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x, a} \left[\omega^{\text{Tr } a(c f(x) - \langle \alpha, x \rangle)} \overline{\psi_k(a)} \right] |c\rangle \langle c|. \end{aligned} \quad (2.59)$$

Since $c \in \mathbb{F}_q^\times$ is invertible, we perform a change of variables. Let $\alpha = c\beta$ for a fresh $\beta \in \mathbb{F}_q^n$. Because c is fixed, this is a bijection on $\mathbb{F}_q^n$. We also substitute $b = ca$, which is a bijection on $\mathbb{F}_q^\times$. This means $a = bc^{-1}$, and by the multiplicity of Dirichlet characters, $\overline{\psi_k(a)} = \overline{\psi_k(bc^{-1})} = \overline{\psi_k(b)} \psi_k(c)$.

Substituting these into the expectation isolates the c -dependence:

$$\begin{aligned} [\hat{F}_f(\chi_{c\beta, k})]_{cc} &= \mathbb{E}_{x, b} \left[\omega^{\text{Tr } b(f(x) - \langle \beta, x \rangle)} \overline{\psi_k(b)} \psi_k(c) \right] \\ &= \psi_k(c) \nu_k(\beta), \end{aligned} \quad (2.60)$$

where $\nu_k(\beta) := \mathbb{E}_{x, b} \left[\omega^{\text{Tr } b(f(x) - \langle \beta, x \rangle)} \overline{\psi_k(b)} \right]$ is a c -independent coefficient. $\square$

Lemma 2.5.11 (Gowers-Hatami to linearity (Convex character Fourier hybrid)). *Suppose $F_f : G \rightarrow U(\mathbb{C}^{q-1})$ defined in Equation 2.36 via $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ is an (ϵ, σ) approximate unitary representation of $(G, \star)$ defined in Equation 2.30. Then, there exists a linear function $\ell : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ such that the relative Hamming distance $\delta(f, \ell)$ is bounded by:*

$$\delta(f, \ell) \leq \frac{q-1}{q} \epsilon. \quad (2.61)$$

Proof. By Gowers-Hatami theorem, we have an isometry $V : \mathcal{H} \rightarrow \mathcal{H}'$ and an exact representation $R : G \rightarrow \text{End}(\mathcal{H}')$, where the Hilbert space $\mathcal{H} \cong \mathbb{C}^{q-1}$ and the space R acts on $\mathcal{H}' = L(G, \mathcal{H})$

$$\mathbb{E}_{g \in G} \Re \langle F_f(g), V^* R(g) V \rangle_\sigma \geq 1 - \epsilon. \quad (2.62)$$

Let $\{|c\rangle\}_{c \in \mathbb{F}_q^\times}$ be the orthonormal basis for $\mathcal{H}$, and the set of characters $\{\chi\}_{\chi \in \hat{G}}$ be an orthonormal basis for $L(G, \mathbb{C})$.

Because $F_f(g)$ is diagonal in $\{|c\rangle\}$ with entries $\omega^{\text{Tr } caf(x)}$, we can rewrite the inner product trace:

$$\langle F_f(g), V^* R(g) V \rangle_\sigma = \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \langle c | F_f^*(g) V^* R(g) V | c \rangle = \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \omega^{-\text{Tr } caf(x)} \langle v_c, R(g) v_c \rangle_{L(G, \mathcal{H})},$$

where $|v_c\rangle = V|c\rangle \in \mathcal{H}'$. Taking the expectation over $g = (x, a)$ and taking the real part, condition Equation 2.62 becomes:

$$\frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x,a} \Re [\omega^{-\text{Tr } caf(x)} \langle v_c, R(x, a) v_c \rangle] \geq 1 - \epsilon. \quad (2.63)$$

A basis for $L(G, \mathcal{H})$ is given by $\{\chi|c\rangle\}_{\chi \in \hat{G}, c \in \mathbb{F}_q^\times}$. By Proposition 2.5.6, the characters are given by

$$\hat{G} = \{\chi_{\alpha,k}\}_{\alpha \in X, k \in [0, \dots, q-2]}, \chi_{\alpha,k}(x, a) = \omega^{\text{Tr}(\alpha, ax)} \psi_k(a).$$

Moreover, the Fourier coefficients of F_f are

$$\begin{aligned} \hat{F}_f(\chi_{\alpha,k}) &= \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x,a} \omega^{\text{Tr } caf(x)} \bar{\chi}_{\alpha,k}(x, a), \\ &= \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x,a} [\omega^{\text{Tr } a(c f(x) - \langle \alpha, x \rangle)} \overline{\psi_k(a)}] |c\rangle \langle c|. \end{aligned} \quad (2.64)$$

The regular representation $R(g)$ acts on $L(G, \mathcal{H})$ so it can be decomposed as

$$R(g) = \sum_{\alpha,k} \chi_{\alpha,k}(g) P_{\alpha,k}, \quad (2.65)$$

where $P_{\alpha,k} = |\chi_{\alpha,k}\rangle \langle \chi_{\alpha,k}| \otimes \mathbb{I}_{\mathcal{H}}$ projects onto the character function subspace (this leaves $\mathcal{H}$ invariant).

Substituting this into the matrix element for v_c , we obtain:

$$\langle v_c, R(g) v_c \rangle = \sum_{\alpha,k} \chi_{\alpha,k}(g) \langle v_c, P_{\alpha,k} v_c \rangle.$$

Let $p_{c,\alpha,k} := \langle v_c, P_{\alpha,k} v_c \rangle$. As $P_{\alpha,k}$ is a projector, then

$$p_{c,\alpha,k} = \langle v_c | \chi_{\alpha,k} \rangle \langle \chi_{\alpha,k} | \otimes \mathbb{I}_{\mathcal{H}} | v_c \rangle \geq 0. \quad (2.66)$$

Let the isometry be given by $V|c\rangle = [F_f(g)]_c$, then the matrix element is given by the Fourier coefficients

$$p_{c,\alpha,k} = \left| [\hat{F}_f(\chi_{\alpha,k})]_{cc} \right|^2. \quad (2.67)$$

By Proposition 2.5.10, let $\alpha = c\beta$, we can show

$$p_{c,c\beta,k} = |\psi_k(c) \nu_k(\beta)|^2 = |\nu_k(\beta)|^2, \quad (2.68)$$

where

$$\nu_k(\beta) := \mathbb{E}_{x,b} [\omega^{\text{Tr } b(f(x) - \langle \beta, x \rangle)} \overline{\psi_k(b)}],$$

is defined in Equation 2.58.

Since V is an isometry $V^*V = \mathbb{I}_{\mathcal{H}}$, we have $\langle c|V^*V|c\rangle = \langle c,c\rangle_{\mathcal{H}} = 1$. Equivalently, we have $\langle v_c, v_c\rangle_{L(G, \mathcal{H})} = 1$. Since $\sum P_{\alpha,k} = I_{L(G, H)}$, it follows that $p_{c,\alpha,k}$ is a valid probability distribution over the characters,

$$\sum_{\alpha,k} p_{c,\alpha,k} = 1, \quad \forall c \in \mathbb{F}_q^\times. \quad (2.69)$$

Equivalently, we have the relabeled probability

$$\sum_{\beta,k} p_{c,c\beta,k} = 1, \quad \forall c \in \mathbb{F}_q^\times. \quad (2.70)$$

Therefore $|\nu_k(\beta)|^2$ is a probability distribution indexed by (k, β) . Because the $k = 0$ probability is a subset, we have

$$\sum_{\beta \in \mathbb{F}_q^n} |\nu_0(\beta)|^2 \leq \sum_{\beta \in \mathbb{F}_q^n} \sum_k |\nu_k(\beta)|^2 = 1. \quad (2.71)$$

Substituting the spectral decomposition Equation 2.65 back into Equation 2.63, we define:

$$\begin{aligned} E &:= \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \mathbb{E}_{x,a} \Re \left[\omega^{-\text{Tr } c a f(x)} \langle v_c, R(x, a) v_c \rangle \right], \\ &= \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \sum_{\alpha,k} p_{c,\alpha,k} \Re \left(\mathbb{E}_{x,a} \left[\omega^{-\text{Tr } c a f(x)} \omega^{\text{Tr } \langle \alpha, a x \rangle} \psi_k(a) \right] \right). \end{aligned}$$

Relabel the sum by $\alpha = c\beta$ and $b = ca$, we have

$$E = \frac{1}{q-1} \sum_{c \in \mathbb{F}_q^\times} \sum_{\beta,k} |\nu_k(\beta)|^2 \Re \left(\mathbb{E}_{x,b} \left[\omega^{\text{Tr } b(\langle \beta, x \rangle - f(x))} \psi_k(c^{-1}) \psi_k(b) \right] \right).$$

Use the indicator function of $\mathbb{F}_q^\times$:

$$\sum_{c \in \mathbb{F}_q^\times} \psi_k(c^{-1}) = (q-1) \mathbb{I}(k=0)$$

we have

$$E = \sum_{\beta} |\nu_0(\beta)|^2 \Re \left(\mathbb{E}_{x,b} \left[\omega^{\text{Tr } b(\langle \beta, x \rangle - f(x))} \right] \right).$$

This is a convex combination (a weighted average) of real values less than 1 with the weight sum to less than 1 by Equation 2.71. Because the average is at least $E \geq 1 - \epsilon$, there must exist at least one specific configuration β^* such that its individual value satisfies:

$$E^* := \Re \left(\mathbb{E}_{x,b} \left[\omega^{\text{Tr } b(\langle \beta^*, x \rangle - f(x))} \right] \right) \geq 1 - \epsilon. \quad (2.72)$$

Let $\ell_\beta(x) := \langle \beta^*, x \rangle$, the inner expectation over $b \in \mathbb{F}_q^\times$ for a fixed x :

$$\mathbb{E}_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b z(x)} = \begin{cases} 1, & z(x) = 0, \\ -\frac{1}{q-1}, & z(x) \neq 0, \end{cases} \quad (2.73)$$

where in the second case we have again used average of roots of unity is $\sum_{b \in \mathbb{F}_q^\times} \omega^{\text{Tr } b} + 1 = 0$.

Evaluating the outer expectation over x and recall the definition of hamming distance

$$\delta(f, \ell_\alpha) := \Pr_x[f(x) \neq \ell_\alpha(x)], \quad (2.74)$$

we find:

$$E^* = 1 \cdot (1 - \delta(f, \ell_\beta)) + \left(\frac{-1}{q-1}\right) \cdot \delta(f, \ell_\beta) = 1 - \delta(f, \ell_\beta) \left(1 + \frac{1}{q-1}\right) = 1 - \delta(f, \ell_\beta) \left(\frac{q}{q-1}\right).$$

Applying our bound from Equation 2.72, we have:

$$1 - \delta(f, \ell_\beta) \left(\frac{q}{q-1}\right) \geq 1 - \epsilon.$$

Rearranging yields the final result: $\exists \ell_\beta = \langle \beta^*, x \rangle$ such that

$$\delta(f, \ell_\beta) \leq \frac{q-1}{q} \epsilon.$$

□

Finally, we are ready to prove the main theorem of BLR linearity test soundness.

Theorem 2.5.1 (Soundness of the BLR test). *Let q be prime. For every $f : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$, $\delta(f, \mathcal{LIN}) \leq \Pr_{a,b,x,y}[af(x) + bf(y) \neq f(ax + by)]$.*

Proof. First denote

$$X := \mathbb{F}_q^n, \quad A := \mathbb{F}_q^\times, \quad \omega := e^{2\pi i/p}.$$

Define a group G by the set $X \times A$ with the binary operation

$$(x, a) \star (y, b) := (b^{-1}x + a^{-1}y, ab). \quad (2.75)$$

Let F_f be a unitary matrix valued lift $F_f : G \rightarrow U(\mathbb{C}^{q-1})$ by the diagonal matrix

$$F_f(x, a) = D(af(x)) := \sum_{c \in A} \omega^{\text{Tr } caf(x)} |c\rangle \langle c|, \quad (2.76)$$

where $\{|c\rangle\}$ denotes the orthonormal basis vectors of a Hilbert space $\mathcal{H} \cong \mathbb{C}^{q-1}$ and the inner product between operators defined by $\langle A, B \rangle_\sigma = \text{Tr}(A^* B \sigma)$ and choose $\sigma = \mathbb{I}/(q-1)$. By Lemma 2.5.7, since f is rejected with probability ε , then F_f is an $(\frac{q}{q-1}\varepsilon, \frac{\mathbb{I}}{q-1})$ approximate representation of $(G, \star)$.

Then by Lemma 2.5.11, there exists a linear function $\ell : \mathbb{F}_q^n \rightarrow \mathbb{F}_q$ such that the relative Hamming distance $\delta(f, \ell)$ is bounded by:

$$\delta(f, \ell) \leq \frac{q-1}{q} \frac{q}{q-1} \varepsilon = \varepsilon. \quad (2.77)$$

This proves the claim

$$\delta(f, \mathcal{LIN}) := \min_\alpha \delta(f, \ell_\alpha) \leq \delta(f, \ell) \leq \varepsilon.$$

□

Miraculously, the loose factor in $(\epsilon' = \frac{q}{q-1}\varepsilon, \frac{\mathbb{I}}{q-1})$ approximate representation F_f cancels with the tighter factor in

$$\delta(f, \ell) \leq \frac{q-1}{q} \epsilon'.$$

Chapter 3

Exponential-length PCP

Definition 3.0.1. Let $m, n > 0$ be integers, $\mathbb{F}$ be a field, $R = \mathbb{F}[x_1, \dots, x_n]$ be its polynomial ring and R_d denote the subspace of polynomials of degree $\leq d$.

$$\text{QESAT}(\mathbb{F}, n) := \{(p_1, \dots, p_m) \in R_2^m \mid \exists a \in \mathbb{F}^n, \forall i \in [m], p_i(a_1, \dots, a_n) = 0\}.$$

For example, $(x + 1, xy + z) \in \text{QESAT}(\mathbb{F}_2, 3)$. The goal of this chapter is to construct an exponential-length PCP for $\text{QESAT}(\mathbb{F}_2, n)$ (Theorem 3.7.8). Since $\text{QESAT}(\mathbb{F}_2, n)$ is NP-complete, this will imply the existence of exponential-length PCPs for all languages in NP.

3.1 Oracle computations and verifiers

Instead of oracle Turing machines, we will define verifiers using Lean monads, specifically the `OracleComp` free monad from `VCV-io`.

Definition 3.1.1 (Oracle specification). An *oracle specification* defines the type of the queries and outputs that an oracle can handle.

Definition 3.1.2. For this section, the *randomness oracle specification* will be of a monadic function that takes nothing and returns an element from a field $\mathbb{F}$.

Definition 3.1.3 (Query implementation). A *query implementation* is a monadic function whose behavior conforms to a given oracle specification.

Definition 3.1.4. The *uniform sampling query implementation* for a field $\mathbb{F}$ is a (mathematical) query implementation for the randomness oracle specification that samples uniformly from $\mathbb{F}$.

This is the the oracle implementation that will be used to prove theorems about our verifiers, since it is compatible with the more classical notion of probability distributions used in the previous chapters.

Remark 3.1.5 (Implementing verifiers). If we want to actually run the verifiers, we can also use concrete implementations based on the IO monad and system randomness for pseudorandom sampling.

Definition 3.1.6 (Oracle context). An *oracle context* is a bundle consisting of an oracle specification and a query implementation for that specification.

Definition 3.1.7 (Randomness oracle). As *randomness oracle*, we will use the uniform sampling query implementation for the randomness oracle specification.

Definition 3.1.8 (Oracle computation). An *oracle computation* is a monadic computation in the free monad generated by an oracle specification, i.e. a description (a “script”) of a program consisting of a chain of oracle queries (compatible with any query implementation of the specification) and deterministic Lean computations.

Definition 3.1.9 (Generic verifier). A *verifier* is a function that takes an input $x \in \{0, 1\}^*$ and produces an oracle computation (with query access to a randomness oracle) that outputs a boolean. This definition can be extended to include other oracles (typically a proof oracle).

We also need to define verifiers that do not take any inputs (for example the BLR verifier). Since they use the same randomness oracle specification as the normal verifiers, they can compose nicely with them.

Definition 3.1.10 (Unit verifier). A *unit verifier* (or just *verifier* when it is clear from the context) is an oracle computation with query access to a randomness oracle that outputs a boolean. This definition can be extended to include other oracles (typically a proof oracle).

Remark 3.1.11 (Time complexity). The price to pay for using these definitions for verifiers instead of Turing machines is that we have no simple way to enforce the time complexity of the verifier. For that, a possible approach would be to restrict the allowed operations to a strict set of operations instead of general Lean computations, and track the time cost using query complexity.

3.2 PCP and LPCP definitions

Definition 3.2.1. A *PCP verifier* is a verifier with query access to a function $\pi : [\ell(|x|)] \rightarrow \Sigma$.

Definition 3.2.2. A language $L \subseteq \{0, 1\}^*$ is in $\text{PCP}[\varepsilon_c, \varepsilon_s, \Sigma, \ell, q, r]$ if there exists a polynomial-time PCP verifier V such that for every $x \in \{0, 1\}^*$, V makes at most $q(|x|)$ queries to the proof oracle and at most $r(|x|)$ randomness queries, and the following holds:

- Completeness: If $x \in L$, then $\exists \pi \in \Sigma^{\ell(|x|)}$, $\Pr[V^\pi(x) = 1] \geq 1 - \varepsilon_c$.
- Soundness: If $x \notin L$, then $\forall \tilde{\pi} \in \Sigma^{\ell(|x|)}$, $\Pr[V^{\tilde{\pi}}(x) = 1] \leq \varepsilon_s$.

Definition 3.2.3. A *LPCP verifier* is a verifier with query access to a linear function $\langle \pi, \cdot \rangle : \Sigma^{\ell(|x|)} \rightarrow \Sigma$.

Definition 3.2.4. A language $L \subseteq \{0, 1\}^*$ is in $\text{LPCP}[\varepsilon_c, \varepsilon_s, \Sigma, \ell, q, r]$ if there exists a polynomial-time LPCP verifier V such that for every $x \in \{0, 1\}^*$, V makes at most $q(|x|)$ queries to the linear proof oracle and at most $r(|x|)$ randomness queries, and the following holds:

- Completeness: If $x \in L$, then $\exists \pi \in \Sigma^{\ell(|x|)}$, $\Pr[V^{\langle \pi, \cdot \rangle}(x) = 1] \geq 1 - \varepsilon_c$.
- Soundness: If $x \notin L$, then $\forall \tilde{\pi} \in \Sigma^{\ell(|x|)}$, $\Pr[V^{\langle \tilde{\pi}, \cdot \rangle}(x) = 1] \leq \varepsilon_s$.

3.3 Simpler BLR verifier for $\mathbb{F}_2$

Definition 3.3.1. Let $V_{\text{BLR}(\mathbb{F}_2, n)}^f$ be the verifier with oracle access to $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$ defined as follows:

1. Sample $x, y \leftarrow \mathbb{F}_2^n$.
2. Accept if and only if $f(x) + f(y) = f(x + y)$.

Theorem 3.3.2. *Let $n > 0$. For every linear function $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$, we have*

$$\Pr \left[V_{\text{BLR}(\mathbb{F}_2, n)}^f = 1 \right] = 1.$$

Proof. □

Theorem 3.3.3. *Let $n > 0$ and $f : \mathbb{F}_2^n \rightarrow \mathbb{F}_2$.*

$$\Pr \left[V_{\text{BLR}(\mathbb{F}_2, n)}^f = 1 \right] \leq 1 - \delta(f, \mathcal{LIN}).$$

Proof. □

3.4 Linear PCP for linear equations

Definition 3.4.1. For a field $\mathbb{F}$ and parameters $m, n \in \mathbb{N}$, we define the language

$$\text{LINEQ}(\mathbb{F}, m, n) := \{(M, c) \in \mathbb{F}^{m \times n} \times \mathbb{F}^m \mid \exists b \in \mathbb{F}^n, Mb = c\}.$$

Definition 3.4.2. Let $V_{\text{LINEQ}(\mathbb{F}, m, n)}^{(b, \cdot)}(M, c)$ be the LPCP verifier defined as follows:

1. Sample $r \leftarrow \mathbb{F}^m$.
2. Query $b \in \mathbb{F}^n$ at $u := M^\top r \in \mathbb{F}^n$.
3. Accept if and only if $\langle b, u \rangle = \langle c, r \rangle$.

Theorem 3.4.3.

$$\text{LINEQ}(\mathbb{F}, m, n) \in \text{LPCP}[\varepsilon_c = 0, \varepsilon_s = 1/|\mathbb{F}|, \Sigma = \mathbb{F}, \ell = |x|, q = 1, r = m].$$

Proof. Consider the LPCP verifier $V_{\text{LINEQ}(\mathbb{F}, m, n)}^{(b, \cdot)}(M, c)$.

Completeness: If $Mb = c$, then for every $r \in \mathbb{F}^m$, we have

$$\langle b, u \rangle = b^\top (M^\top r) = (Mb)^\top r = \langle c, r \rangle.$$

Soundness: If $Mb \neq c$, then

$$\begin{aligned} \Pr_{r \leftarrow \mathbb{F}^m} [\langle b, u \rangle = \langle c, r \rangle] &= \Pr_{r \leftarrow \mathbb{F}^m} [\langle Mb, r \rangle = \langle c, r \rangle] \\ &= \Pr_{r \leftarrow \mathbb{F}^m} \left[\sum_{i=1}^m (Mb - c)_i \cdot r_i = 0 \right] \leq 1/|\mathbb{F}|, \end{aligned}$$

where the last inequality follows from the Schwartz-Zippel lemma. □

3.5 Linear PCP for tensor checks

Definition 3.5.1. For a field $\mathbb{F}$, vectors $a \in \mathbb{F}^n$ and $b \in \mathbb{F}^m$, the tensor product $a \otimes b \in \mathbb{F}^{n \times m}$ is the matrix defined by

$$(a \otimes b)_{i,j} := a_i b_j \quad \forall i \in [n], j \in [m].$$

We write $\text{flat}(a \otimes b) \in \mathbb{F}^{nm}$ for the row-concatenation of $a \otimes b$.

For our purposes, we want a test that checks whether $\text{flat}(a \otimes a) = b$

Definition 3.5.2. For a field $\mathbb{F}$ and $n \in \mathbb{N}$, define the tensor test language

$$\text{TENSORQ}(\mathbb{F}, n) := \{(a, b) \in \mathbb{F}^n \times \mathbb{F}^{n^2} \mid b = \text{flat}(a \otimes a)\}.$$

Definition 3.5.3. Given $(a, b) \in \mathbb{F}^n \times \mathbb{F}^{n^2}$, the LPCP verifier $V_{\text{TENSORQ}(\mathbb{F}, n)}^{\langle \pi, \cdot \rangle}(a, b)$ uses a proof $\pi = \alpha \parallel \beta \in \mathbb{F}^{n+n^2}$, intended to be the linear encoding of the public input (a, b) . Since the proof may be dishonest, the verifier checks consistency with (a, b) on the sampled queries before checking the tensor relation:

1. Sample $s, t \leftarrow \mathbb{F}^n$.
2. Query the α -part of π at s and at t , obtaining y_s and y_t .
3. Query the β -part of π at $\text{flat}(s \otimes t)$, obtaining $y_{s,t}$.
4. Accept if and only if

$$y_s = \langle a, s \rangle, \quad y_t = \langle a, t \rangle, \quad y_{s,t} = \langle b, \text{flat}(s \otimes t) \rangle, \quad y_{s,t} = y_s y_t.$$

Definition 3.5.4. In settings where a and b are supplied only by the proof oracle, the tensor self-check uses the same three queries but has no public input to compare against. Given oracle access to $\pi = \alpha \parallel \beta \in \mathbb{F}^{n+n^2}$, sample $s, t \leftarrow \mathbb{F}^n$, query

$$y_s = \langle \alpha, s \rangle, \quad y_t = \langle \alpha, t \rangle, \quad y_{s,t} = \langle \beta, \text{flat}(s \otimes t) \rangle,$$

and accept if and only if $y_{s,t} = y_s y_t$. This is the tensor component used inside the QESAT verifier.

Lemma 3.5.5. For every $a, s, t \in \mathbb{F}^n$,

$$\langle \text{flat}(a \otimes a), \text{flat}(s \otimes t) \rangle = \langle a, s \rangle \langle a, t \rangle.$$

Proof.

$$\langle \text{flat}(a \otimes a), \text{flat}(s \otimes t) \rangle = \sum_{i,j} a_i a_j s_i t_j = \left(\sum_i a_i s_i \right) \left(\sum_j a_j t_j \right) = \langle a, s \rangle \langle a, t \rangle.$$

□

Lemma 3.5.6. Let $\mathbb{F}$ be a finite field. If $b \neq \text{flat}(a \otimes a)$, then for every proof $\pi \in \mathbb{F}^{n+n^2}$,

$$\Pr_{s,t} \left[V_{\text{TENSORQ}(\mathbb{F}, n)}^{\langle \pi, \cdot \rangle}(a, b) = 1 \right] \leq \frac{2|\mathbb{F}| - 1}{|\mathbb{F}|^2}.$$

Proof. Suppose $b \neq \text{flat}(a \otimes a)$, i.e., there exists $i^*, j^* \in [n]$ such that $b_{i^*j^*} \neq a_{i^*} \cdot a_{j^*}$. For each $i \in [n]$ define the linear polynomial $p_i(t) := \sum_{j \in [n]} (b_{ij} - a_i a_j) t_j$. For fixed s, t , if the verifier accepts, then its consistency checks force $y_s = \langle a, s \rangle$, $y_t = \langle a, t \rangle$, and $y_{s,t} = \langle b, \text{flat}(s \otimes t) \rangle$. Thus acceptance can only occur when

$$\langle b, \text{flat}(s \otimes t) \rangle - \langle a, s \rangle \langle a, t \rangle = \sum_{i \in [n]} \sum_{j \in [n]} (b_{ij} - a_i a_j) s_i t_j = \sum_{i \in [n]} p_i(t) s_i.$$

Immediate application of Polynomial Identity Testing yields a lower bound of $1 - \frac{2}{|\mathbb{F}|}$ but we can do better. Since $b_{i^*j^*} \neq a_{i^*} a_{j^*}$, the polynomial p_{i^*} is nonzero, so by the Schwartz–Zippel lemma applied to t ,

$$\Pr_t[p_{i^*}(t) \neq 0] \geq 1 - \frac{1}{|\mathbb{F}|}.$$

Conditioned on $p_{i^*}(t) \neq 0$, the expression $\sum_i p_i(t) s_i$ is a nonzero linear polynomial in s , so by the Schwartz–Zippel lemma applied to s ,

$$\Pr_{s \leftarrow \mathbb{F}^n} \left[\sum_i p_i(t) s_i \neq 0 \mid p_{i^*}(t) \neq 0 \right] \geq 1 - \frac{1}{|\mathbb{F}|}.$$

Hence

$$\Pr_{s,t} [V_{\text{TensorQ}(\mathbb{F},n)}^{(\pi,\cdot)}(a,b) = 0] \geq \left(1 - \frac{1}{|\mathbb{F}|}\right)^2,$$

and therefore

$$\Pr_{s,t} [V_{\text{TensorQ}(\mathbb{F},n)}^{(\pi,\cdot)}(a,b) = 1] \leq 1 - \left(1 - \frac{1}{|\mathbb{F}|}\right)^2 = \frac{2|\mathbb{F}| - 1}{|\mathbb{F}|^2}.$$

□

Lemma 3.5.7. *Let $\mathbb{F}$ be a finite field and let $\pi = \alpha \parallel \beta \in \mathbb{F}^{n+n^2}$. If $(\alpha, \beta) \notin \text{TensorQ}(\mathbb{F}, n)$, then the tensor self-check accepts with probability at most*

$$\frac{2|\mathbb{F}| - 1}{|\mathbb{F}|^2}.$$

Proof. Apply 3.5.6 to the public pair (α, β) . The self-check is exactly the multiplicativity part of that verifier, with (α, β) read from the proof oracle itself. □

Theorem 3.5.8. *For every finite field $\mathbb{F}$ and $n \in \mathbb{N}$,*

$$\text{TensorQ}(\mathbb{F}, n) \in \text{LPCP} \left[\varepsilon_c = 0, \varepsilon_s = \frac{2|\mathbb{F}| - 1}{|\mathbb{F}|^2}, \Sigma = \mathbb{F}, \ell = n + n^2, q = 3, r = 2n \right].$$

Proof. Use the verifier of 3.5.3. If $(a, b) \in \text{TensorQ}(\mathbb{F}, n)$, the honest proof $\pi = (a, b)$ passes the three input-consistency checks identically, and the multiplicativity check is 3.5.5. Soundness is 3.5.6. The verifier makes three proof queries and uses two random vectors in $\mathbb{F}^n$, which gives the displayed query and randomness bounds. □

Corollary 3.5.9.

$$\text{TensorQ}(\mathbb{F}_2, n) \in \text{LPCP} [\varepsilon_c = 0, \varepsilon_s = 3/4, \Sigma = \mathbb{F}_2, \ell = n + n^2, q = 3, r = 2n].$$

Proof. Specialize 3.5.8 to $\mathbb{F}_2$, where $|\mathbb{F}_2| = 2$. □

3.6 Linear PCP to PCP

Lemma 3.6.1 (Chernoff bound). *Let $X_1, \dots, X_N$ be independent random variables over $\{0, 1\}$, and let $\mu = \mathbb{E}[\sum_{i=1}^N X_i]$. For any $\Delta > 0$ one has*

$$\Pr\left[\sum_{i=1}^N X_i \geq \mu + \Delta\right] \leq \exp(-2\Delta^2/N)$$

Proof. Write $p_i = \mathbb{E}[X_i]$, so that $\mu = \sum_{i=1}^N p_i$. We first prove the elementary exponential-moment estimate

$$\mathbb{E}[\exp(\lambda(X_i - p_i))] \leq \exp(\lambda^2/8)$$

for every $\lambda > 0$ and every i . Since X_i is $\{0, 1\}$ -valued,

$$\mathbb{E}[\exp(\lambda(X_i - p_i))] = (1 - p_i) \exp(-\lambda p_i) + p_i \exp(\lambda(1 - p_i)).$$

Equivalently, it is enough to show

$$\log(1 - p_i + p_i \exp(\lambda)) - p_i \lambda \leq \lambda^2/8.$$

For $p \in [0, 1]$, set $f(\lambda) = \log(1 - p + p \exp(\lambda)) - p\lambda$. Then $f(0) = f'(0) = 0$, and

$$f''(\lambda) = \frac{p \exp(\lambda)}{1 - p + p \exp(\lambda)} \left(1 - \frac{p \exp(\lambda)}{1 - p + p \exp(\lambda)}\right) \leq \frac{1}{4}.$$

Hence, for $\lambda > 0$,

$$f(\lambda) = \int_0^\lambda \int_0^u f''(v) dv du \leq \int_0^\lambda \int_0^u \frac{1}{4} dv du = \lambda^2/8.$$

Let $S = \sum_{i=1}^N X_i$. By independence and the estimate above,

$$\mathbb{E}[\exp(\lambda(S - \mu))] = \prod_{i=1}^N \mathbb{E}[\exp(\lambda(X_i - p_i))] \leq \exp(N\lambda^2/8).$$

Markov's inequality gives, for every $\lambda > 0$,

$$\Pr[S \geq \mu + \Delta] = \Pr[\exp(\lambda(S - \mu)) \geq \exp(\lambda\Delta)] \leq \exp(-\lambda\Delta) \mathbb{E}[\exp(\lambda(S - \mu))]$$

Using the moment bound above, this is at most

$$\exp(-\lambda\Delta + N\lambda^2/8).$$

Taking $\lambda = 4\Delta/N$ gives

$$\Pr[S \geq \mu + \Delta] \leq \exp(-2\Delta^2/N).$$

□

Definition 3.6.2. For an LPCP verifier making at most q linear queries, set

$$t := 8(\lceil \log_2(100q) \rceil + 1).$$

This is the concrete logarithmic shift count used in Lean. The converted PCP verifier uses t independent random shifts to decode each linear query by plurality.

Definition 3.6.3 (Encoded proof). For a linear LPCP proof $\pi \in \mathbb{F}_2^\ell$, its PCP encoding $\hat{\pi}_{\text{lin}} \in \mathbb{F}_2^\ell$ is the truth table

$$\hat{\pi}_{\text{lin}}(a) := \langle \pi, a \rangle, \quad a \in \mathbb{F}_2^\ell.$$

Definition 3.6.4. Given a PCP proof $\hat{\pi} : \mathbb{F}_2^\ell \rightarrow \mathbb{F}_2$, shifts $r_1, \dots, r_t \in \mathbb{F}_2^\ell$, and a linear query $a \in \mathbb{F}_2^\ell$, define the locally corrected answer

$$\text{Corr}_{\hat{\pi}}(a; r_1, \dots, r_t) := \text{plurality}_{i \in [t]}(\hat{\pi}(a + r_i) - \hat{\pi}(r_i)).$$

Definition 3.6.5. Let V be an LPCP verifier over $\mathbb{F}_2$. The converted PCP verifier $\bar{V}$, with oracle access to $\hat{\pi} : \mathbb{F}_2^\ell \rightarrow \mathbb{F}_2$, proceeds as follows on input x :

1. If $V_{\text{BLR}(\mathbb{F}_2, n)}^{\hat{\pi}} = 0$, reject. Otherwise, proceed to the next step.
2. Sample t shifts $r_1, \dots, r_t \in \mathbb{F}_2^\ell$, with t as in ??.
3. Simulate $V(x)$. Whenever V asks a linear query $a \in \mathbb{F}_2^\ell$, answer it by $\text{Corr}_{\hat{\pi}}(a; r_1, \dots, r_t)$.

Lemma 3.6.6 (Completeness). *If $\Pr[V^{\langle \pi, \cdot \rangle}(x) = 1] \geq 1 - \varepsilon_c$, then $\Pr[\bar{V}^{\hat{\pi}_{\text{lin}}}(x) = 1] \geq 1 - \varepsilon_c$.*

Proof. The encoded proof is linear, so the BLR test accepts with probability 1 by 3.3.2. For every query a and every shift r_i ,

$$\hat{\pi}_{\text{lin}}(a + r_i) - \hat{\pi}_{\text{lin}}(r_i) = \langle \pi, a \rangle.$$

Thus every plurality answer equals the corresponding LPCP oracle answer, and the simulation has exactly the same acceptance probability as V . $\square$

Lemma 3.6.7 (Soundness, far case). *If $\delta(\hat{\pi}, \mathcal{LJN}) \geq 1/8$, then*

$$\Pr[\bar{V}^{\hat{\pi}}(x) = 1] \leq 7/8.$$

Proof. By 3.3.3, the BLR precheck accepts with probability at most $1 - \delta(\hat{\pi}, \mathcal{LJN}) \leq 7/8$. Since $\bar{V}$ accepts only when this test accepts, its acceptance probability is at most $7/8$. $\square$

Lemma 3.6.8. *If $\delta(\hat{\pi}, f) < 1/8$ for some linear function f of the form $u \mapsto \langle \pi, u \rangle$, then*

$$\Pr_{r_1, \dots, r_t} [\text{Corr}_{\hat{\pi}}(a; r_1, \dots, r_t) \neq \langle \pi, a \rangle] \leq \exp(-t/8), \quad \forall a \in \mathbb{F}_2^\ell.$$

Proof. For one shift r , the corrected sample can be wrong only if $\hat{\pi}(a + r)$ or $\hat{\pi}(r)$ disagrees f . A union bound gives error probability at most $2 \cdot (1/8) = 1/4$. If X_i is the indicator of failure for shift r_i , then $\mathbb{E}[\sum_i X_i] \leq t/4$. The plurality answer is wrong only if $\sum_i X_i \geq t/2$, so ?? with $\Delta = t/4$ gives $\exp(-t/8)$. $\square$

Lemma 3.6.9. *Assume every possible linear query a is locally corrected incorrectly with probability at most η . If the simulated LPCP computation makes at most q proof queries, then its acceptance probability under the corrected PCP oracle is at most the acceptance probability under the nearby linear proof plus $q\eta$.*

Proof. Expose the LPCP computation one oracle query at a time. At a proof query, the two simulations can diverge only when the local correction for the queried linear form is wrong; by assumption this has probability at most η . Inducting over the query-bound derivation gives an additive loss of at most $q\eta$. $\square$

Lemma 3.6.10 (Soundness, near case). *If $x \notin L$, V has LPCP soundness ε_s , and $\delta(\hat{\pi}, f) < 1/8$ for some linear function f of the form $u \mapsto \langle \pi, u \rangle$, then*

$$\Pr[\bar{V}^{\hat{\pi}}(x) = 1] \leq \varepsilon_s + \frac{1}{100}.$$

Proof. By 3.6.8, every possible linear query has local correction error at most $\eta = \exp(-t/8)$. The hybrid estimate 3.6.9 compares the corrected simulation with the simulation using the nearby linear proof and loses at most $q\eta$. The choice of t in 3.6.2 gives $q\eta \leq 1/100$. The nearby linear proof is a valid LPCP proof, so its acceptance probability is at most ε_s because $x \notin L$. $\square$

Lemma 3.6.11 (Soundness). *For every $x \notin L$ and every PCP proof $\hat{\pi}$,*

$$\Pr[\bar{V}^{\hat{\pi}}(x) = 1] \leq \max\{7/8, \varepsilon_s + 1/100\}.$$

Proof. If $\delta(\hat{\pi}, \mathcal{LJN}) \geq 1/8$, apply 3.6.7. Otherwise, apply 3.6.10. $\square$

Lemma 3.6.12.

$$\begin{aligned} & \text{LPCP}[\varepsilon_c, \varepsilon_s, \Sigma = \mathbb{F}_2, \ell, q, r] \subseteq \\ & \text{PCP}[\varepsilon_c, \varepsilon'_s = \max\{7/8, \varepsilon_s + 1/100\}, \Sigma' = \mathbb{F}_2, \ell' = 2^\ell, q' = 3 + 2tq, r' = r + (2+t)\ell]. \end{aligned}$$

Proof. Let $L \in \text{LPCP}[\varepsilon_c, \varepsilon_s, \Sigma = \mathbb{F}_2, \ell, q, r]$ and let V be the LPCP verifier for L . Use the converted PCP verifier $\bar{V}$ defined in 3.6.5. Completeness is 3.6.6, and soundness is 3.6.11. The BLR test makes three proof queries. Each simulated LPCP query uses two PCP queries for each of the t shifts, so the query bound is $3 + 2tq$. The randomness consists of the original r random bits, 2ℓ bits for the BLR test, and $t\ell$ bits for the shifts. $\square$

3.7 PCP for QESAT

Definition 3.7.1. Let $x = (p_1, \dots, p_m)$ be a QESAT($\mathbb{F}_2, n$) instance. For each quadratic polynomial p_i , form a row $M_i \in \mathbb{F}_2^{n+n^2}$ by replacing every linear monomial X_j by the variable a_j and every quadratic monomial $X_j X_k$ by the tensor variable b_{jk} . Let

$$M_x := \begin{bmatrix} M_1 \\ \vdots \\ M_m \end{bmatrix} \in \mathbb{F}_2^{m \times (n+n^2)} \quad \text{and} \quad c_x := \begin{bmatrix} -\text{const}(p_1) \\ \vdots \\ -\text{const}(p_m) \end{bmatrix} \in \mathbb{F}_2^m,$$

where $\text{const}(p_i)$ is the constant term of p_i . Then $M_x(a \parallel b) = c_x$ expresses the original equations after replacing $a_j a_k$ by b_{jk} .

Definition 3.7.2. Let $x = (p_1, \dots, p_m) \in R_2^m$ be a QESAT instance as defined in 3.0.1. Let $V_{\text{QESAT}(\mathbb{F}_2, n)}^\pi(x)$ be the LPCP verifier with oracle access to $\pi = a \parallel b \in \mathbb{F}_2^{n+n^2}$, defined as follows:

1. Sample $r \leftarrow \mathbb{F}_2^m$ and $s, t \leftarrow \mathbb{F}_2^n$.
2. Construct $(M_x, c_x) \in \mathbb{F}_2^{m \times (n+n^2)} \times \mathbb{F}_2^m$ as in 3.7.1.
3. Check that $V_{\text{LINEQ}(\mathbb{F}_2, m, n+n^2)}^{\langle \pi, \cdot \rangle}(M_x, c_x) = 1$.

4. Query π at $(s, 0)$, $(t, 0)$, and $(0, \text{flat}(s \otimes t))$ and check that

$$\langle b, \text{flat}(s \otimes t) \rangle = \langle a, s \rangle \langle a, t \rangle.$$

Equivalently, this last step is the tensor self-check of 3.5.4 specialized to $\mathbb{F}_2$.

Lemma 3.7.3 (Completeness). *If $x \in \text{QESAT}(\mathbb{F}_2, n)$, then exists a proof $\pi \in \mathbb{F}_2^{n+n^2}$ such that*

$$\Pr[V_{\text{QESAT}(\mathbb{F}_2, n)}^\pi(x) = 1] = 1.$$

Proof. Let $a \in \mathbb{F}_2^n$ satisfy all equations and set $b = \text{flat}(a \otimes a)$. By construction of 3.7.1, $M_x(a \parallel b) = c_x$, so the LINEQ check accepts for every r . The tensor self-check accepts for every s, t by 3.5.5. $\square$

Lemma 3.7.4 (Soundness). *If $x \notin \text{QESAT}(\mathbb{F}_2, n)$, then for every proof $\tilde{\pi} \in \mathbb{F}_2^{n+n^2}$,*

$$\Pr[V_{\text{QESAT}(\mathbb{F}_2, n)}^{\tilde{\pi}}(x) = 1] \leq 3/4.$$

Proof. If some p_i has total degree > 2 , the verifier rejects immediately, so the acceptance probability is 0. Assume now that every p_i is quadratic and write $\pi = a \parallel b$. If $b \neq \text{flat}(a \otimes a)$, then the tensor self-check accepts with probability at most $3/4$ by 3.5.7 specialized to $\mathbb{F}_2$, and the full verifier accepts only if that check accepts.

Otherwise $b = \text{flat}(a \otimes a)$. Since $x \notin \text{QESAT}(\mathbb{F}_2, n)$, the assignment a does not satisfy all equations. By 3.7.1, $M_x(a \parallel b) \neq c_x$. The LINEQ check therefore accepts with probability at most $1/|\mathbb{F}_2| = 1/2$ by ???. Again the full verifier accepts only if this check accepts, so the acceptance probability is at most $1/2 \leq 3/4$. $\square$

Theorem 3.7.5.

$$\text{QESAT}(\mathbb{F}_2, n) \in \text{LPCP}[\varepsilon_c = 0, \varepsilon_s = 3/4, \Sigma = \mathbb{F}_2, \ell = n + n^2, q = 4, r = |x| + 2n].$$

Proof. Use the verifier 3.7.2. It makes one query for the LINEQ check and three queries for the tensor self-check, for a total of four. Its randomness is the LINEQ randomness plus two vectors in $\mathbb{F}_2^n$. Completeness and soundness are 3.7.3 and 3.7.4. $\square$

Theorem 3.7.6.

$$\text{QESAT}(\mathbb{F}_2, n) \in \text{PCP}[\varepsilon_c = 0, \varepsilon_s = 7/8, \Sigma = \mathbb{F}_2, \ell = 2^{|x|}, q = O(1), r = |x|^{O(1)}].$$

Proof. Apply 3.6.12 to 3.7.5. Since the LPCP soundness is $3/4$, the converted soundness is

$$\max\{7/8, 3/4 + 1/100\} = 7/8.$$

The query bound remains constant because $q = 4$, and the randomness is polynomial in the input size. The Lean formalization pads the proof oracle from length 2^{n+n^2} to length $2^{|x|}$. $\square$

Lemma 3.7.7.

$$\text{PCP}[0, 7/8, \mathbb{F}_2, 2^{|x|}, O(1), |x|^{O(1)}] \subseteq \text{PCP}[0, 1/2, \mathbb{F}_2, 2^{|x|}, O(1), |x|^{O(1)}].$$

Proof. Run the $7/8$ -sound verifier independently six times and accept only if all six runs accept. Completeness error remains 0. For $x \notin L$, the acceptance probability is at most $(7/8)^6 \leq 1/2$. Query and randomness bounds are multiplied by the constant 6. $\square$

Theorem 3.7.8.

$$\text{QESAT}(\mathbb{F}_2, n) \in \text{PCP}[\varepsilon_c = 0, \varepsilon_s = 1/2, \Sigma = \mathbb{F}_2, \ell = 2^{|x|}, q = O(1), r = |x|^{O(1)}].$$

Proof. Combine 3.7.6 with the constant repetition step 3.7.7. $\square$

Bibliography

- [Chi24] Alessandro Chiesa. Foundations of probabilistic proofs (online course), 2024. Course slides; Lectures 7 and 8.
- [GH17] W. T. Gowers and Omid Hatami. Inverse and stability theorems for approximate representations of finite groups. *Sbornik: Mathematics*, 208(12):1784–1817, 2017.
- [Mor20] Sophie Morel. Robust self-testing via approximate representations. https://perso.ens-lyon.fr/mikael.de.la.salle/GDT/8-Sophie-approximate_reps.pdf, 2020. Lecture slides from the Connes–Tsirelson working group, 7 May 2020.
- [Vid22] Thomas Vidick. QMATH Summer School on the Mathematics of Entanglement via Nonlocal Games: Course Notes. https://qmath.ku.dk/events/conferences/quantum-entanglement/Vidick_-_course_notes.pdf, 2022. Course notes; includes Theorem 2.3, the Gowers–Hatami theorem.